

# Computer Vision

Local Features

# Image matching



by [Diva Sian](#)



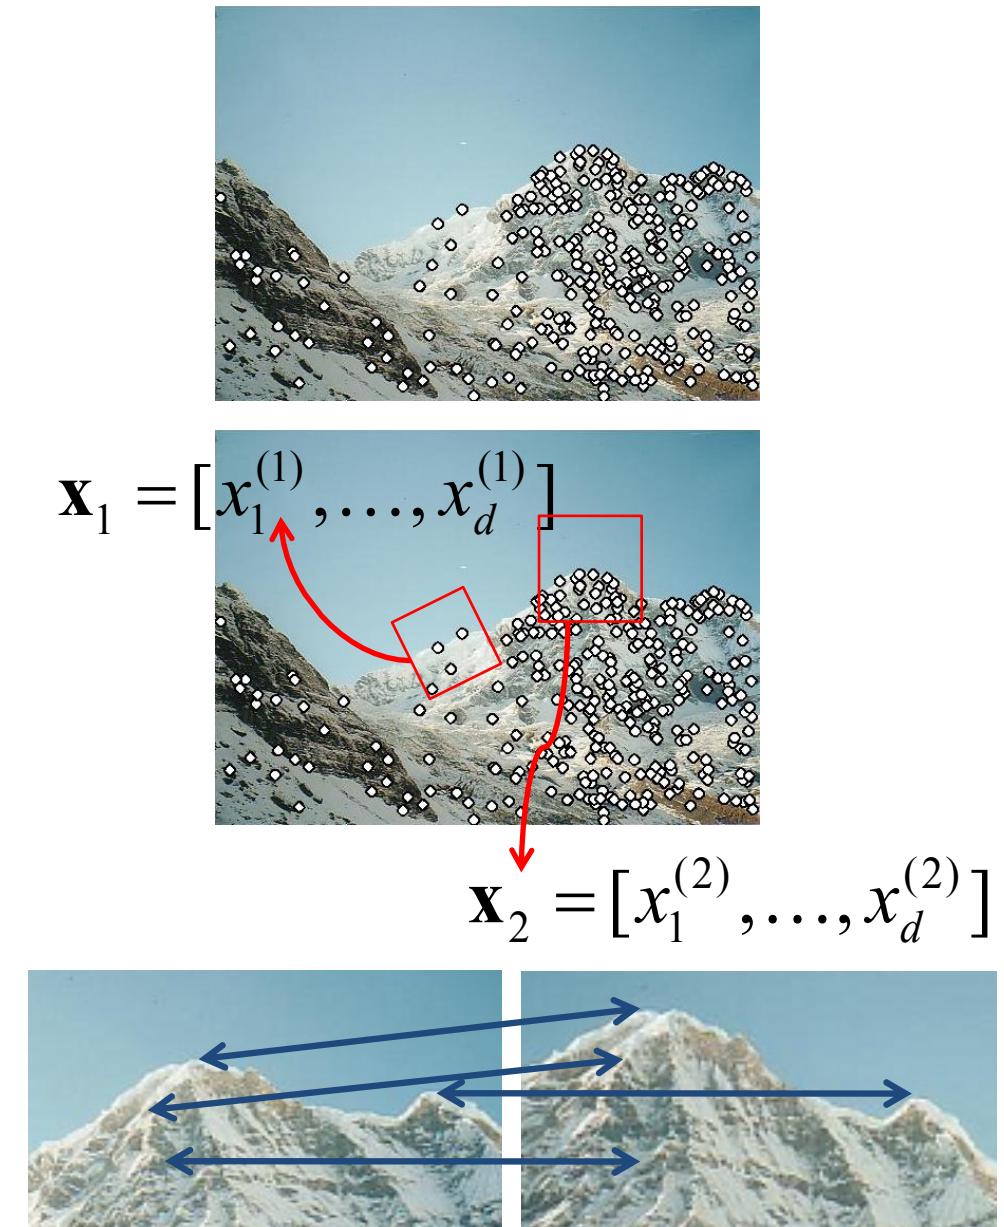
by [swashford](#)





# Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views



# Repeatability

- We want to detect (at least some of) the same points in both images.

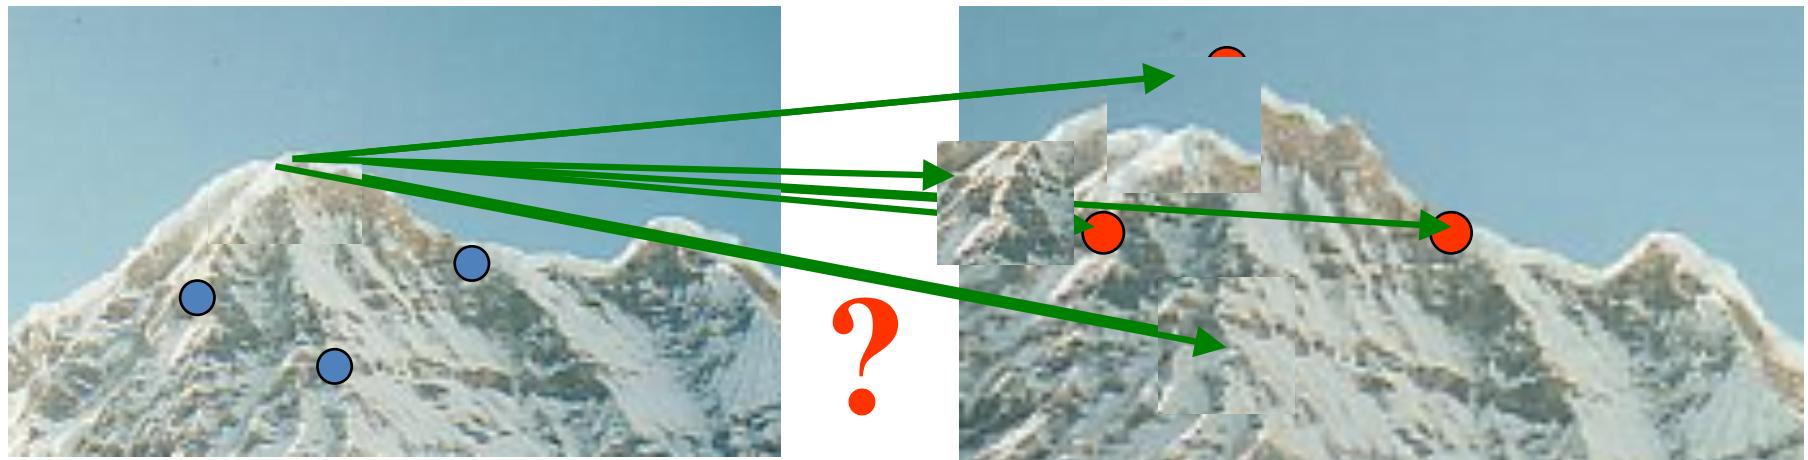


No chance to find true matches!

- Yet we have to be able to run the detection procedure *independently* per image.

# Distinctiveness

- We want to be able to reliably determine which point goes with which.



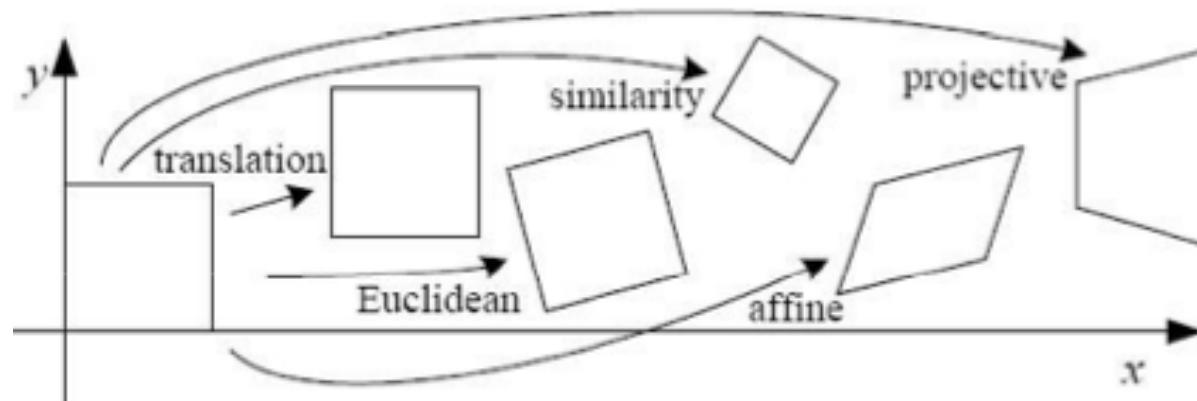
- Must provide some invariance to geometric and photometric differences between the two views.

# Invariant local features

## Geometric transformations

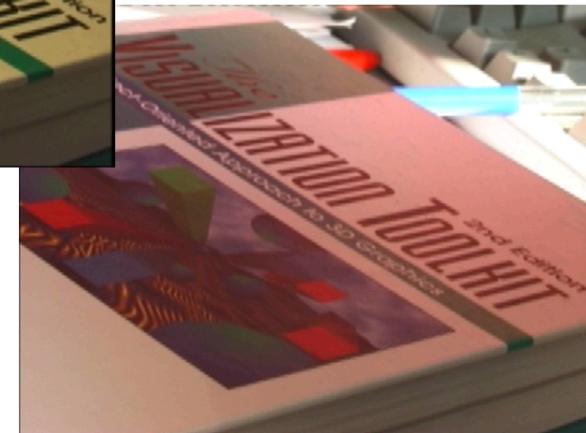
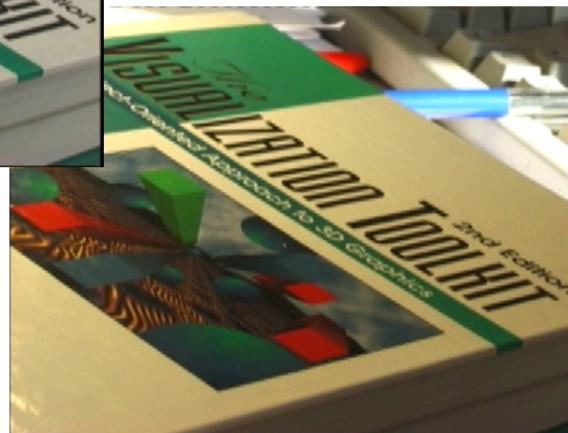
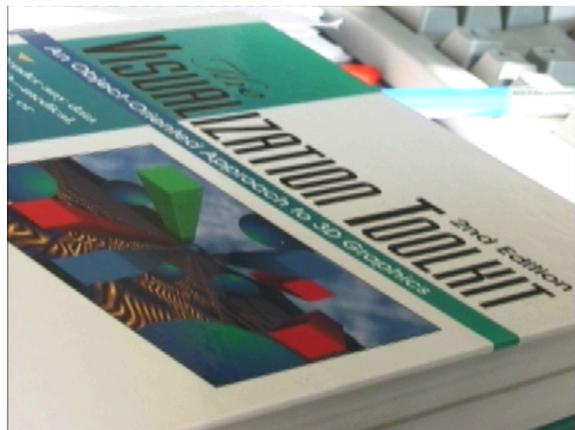
- Translation
- Euclidean (translation + rotation)
- Similarity (translation + rotation + scale)
- Affine transformations
- Projective transformations

**Only holds  
for planar  
patches**



# Invariant local features

## Photometric transformations



- Often modeled as a linear transformation:
  - Scaling + Offset

# Requirements

- Region extraction needs to be **repeatable** and **accurate**
  - Invariant to translation, rotation, scale changes
  - Robust or covariant to out-of-plane (affine) transformations
  - Robust to lighting variations, noise, blur, quantization
- **Locality:** Features are local, therefore robust to occlusion and clutter.
- **Quantity:** We need a sufficient number of regions to cover the object.
- **Distinctiveness:** The regions should contain “interesting” structure.
- **Efficiency:** Close to real-time performance.

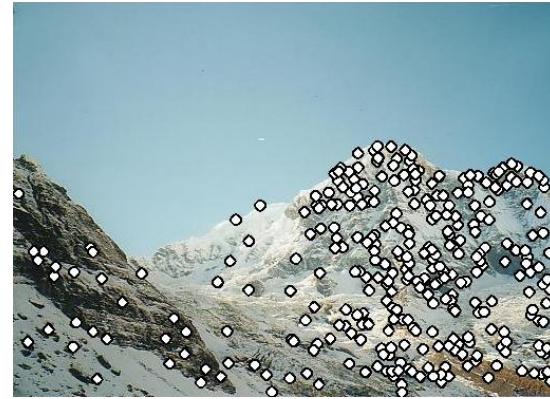
# Many Existing Detectors Available

- Hessian & Harris [Beaudet '78], [Harris '88]
- Laplacian, DoG [Lindeberg '98], [Lowe '99]
- Harris-/Hessian-Laplace [Mikolajczyk & Schmid '01]
- Harris-/Hessian-Affine [Mikolajczyk & Schmid '04]
- EBR and IBR [Tuytelaars & Van Gool '04]
- MSER [Matas '02]
- Salient Regions [Kadir & Brady '01]
- Others...

*Those detectors have become a basic building block for many recent applications in Computer Vision.*

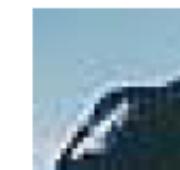
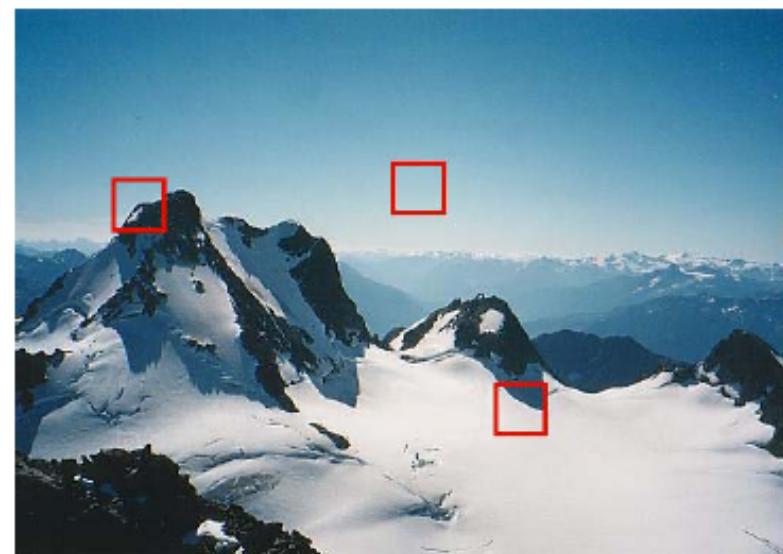
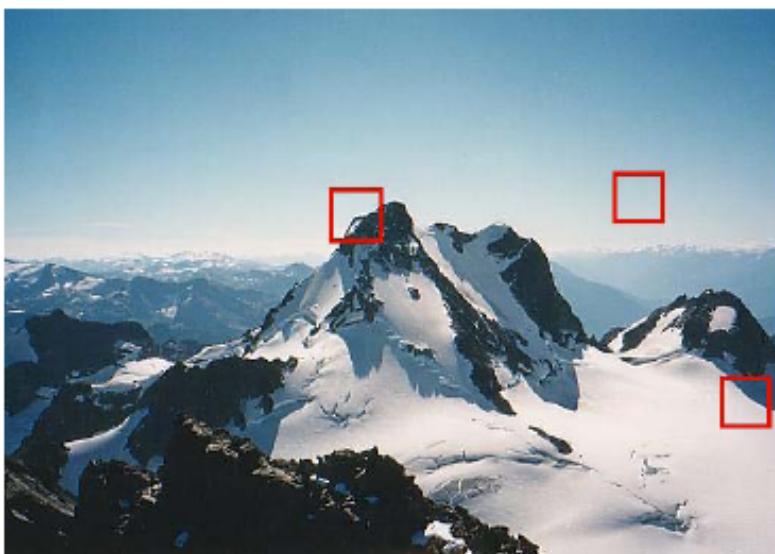
# Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views



# Local features detectors

- What are salient features that can be *detected* in multiple views?



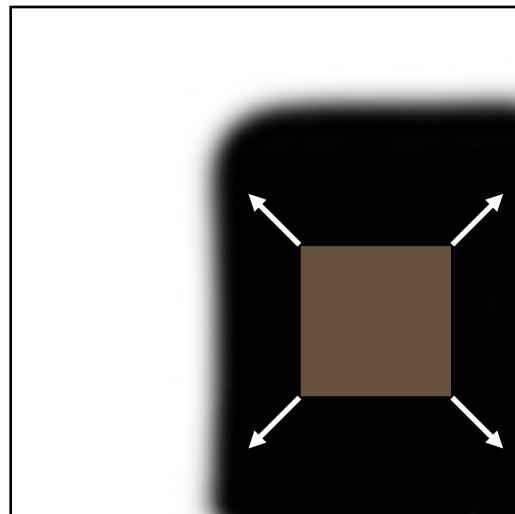
# Local features detectors

- What points would be better?

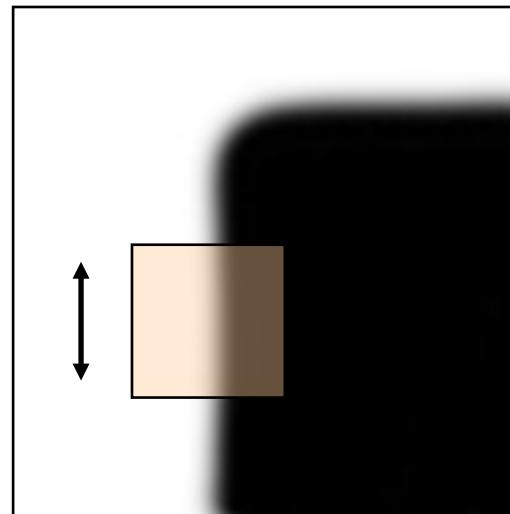


# Corners

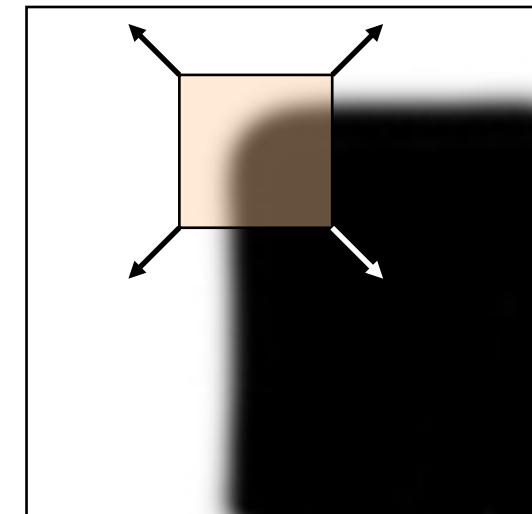
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



**“flat” region:**  
no change in  
all directions



**“edge”:**  
no change  
along the edge  
direction



**“corner”:**  
significant  
change in all  
directions

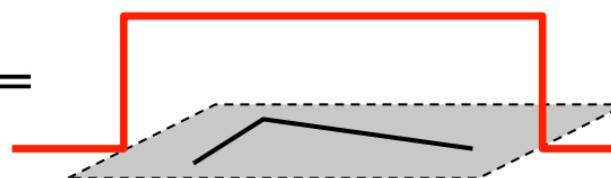
# Harris Detector Formulation

- Change of intensity for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

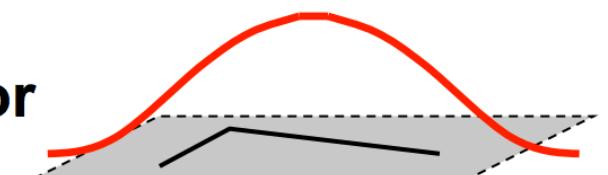
**Window function**      **Shifted intensity**      **Intensity**

**Window function**  $w(x, y) =$



1 in window, 0 outside

or



Gaussian

# Harris Detector Formulation

This measure of change can be approximated by:

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a  $2 \times 2$  matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Gradient with  
respect to x, times  
gradient with  
respect to y

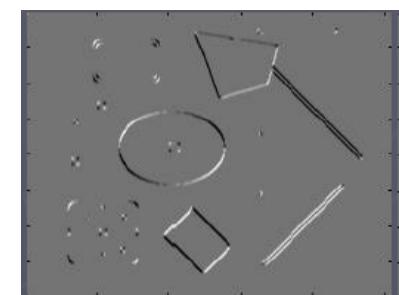
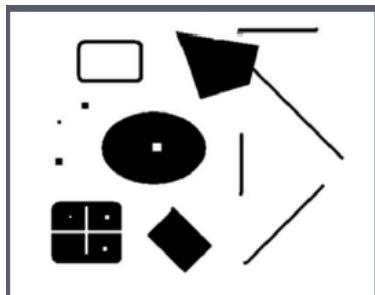
Sum over image region – area  
we are checking for corner

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$

# Harris Detector Formulation

$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point).



**Notation:**

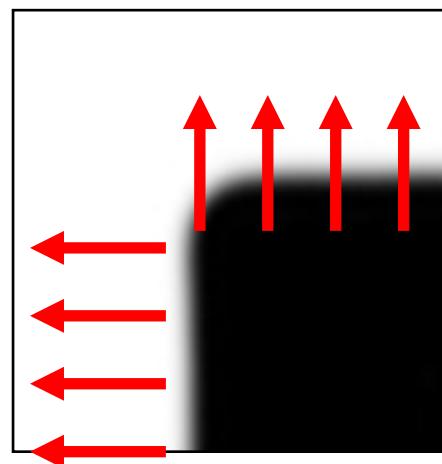
$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# What does this matrix reveal?

First, consider an axis-aligned corner:



# What does this matrix reveal?

First, consider an axis-aligned corner:

$$M = \sum \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

This means dominant gradient directions align with x or y axis

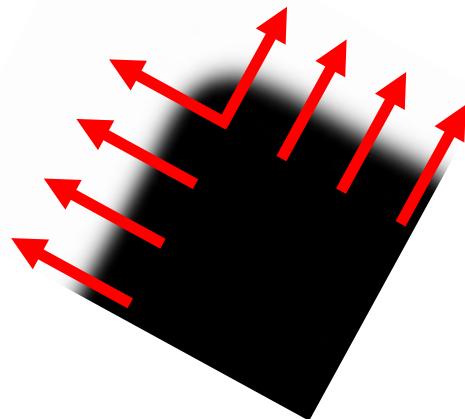
Look for locations where both  $\lambda$ 's are large.

If either  $\lambda$  is close to 0, then this is not corner-like.

What if we have a corner that is not aligned with the image axes?

# What does this matrix reveal?

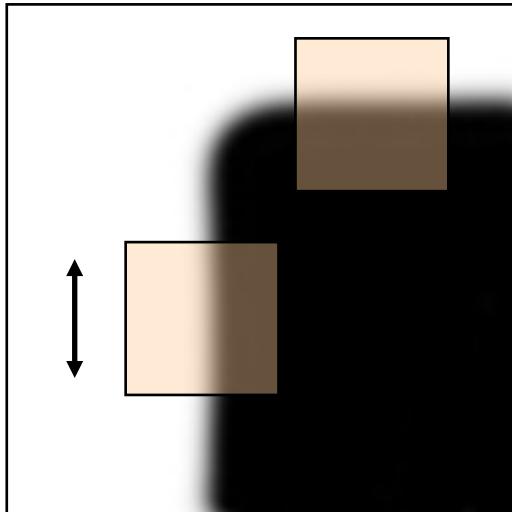
Since  $M$  is symmetric, we have  $M = X \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} X^T$



$$Mx_i = \lambda_i x_i$$

The *eigenvalues* of  $M$  reveal the amount of intensity change in the two principal orthogonal gradient directions in the window.

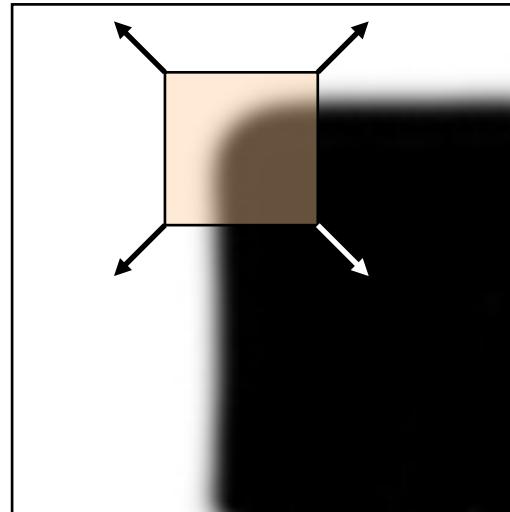
# Corner response function



“edge”:

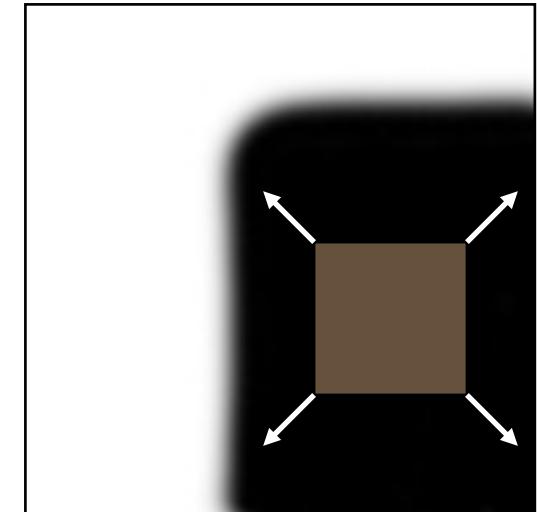
$$\lambda_1 \gg \lambda_2$$

$$\lambda_2 \gg \lambda_1$$



“corner”:

$\lambda_1$  and  $\lambda_2$  are large,  
 $\lambda_1 \sim \lambda_2$ ;



“flat” region

$\lambda_1$  and  $\lambda_2$  are  
small;

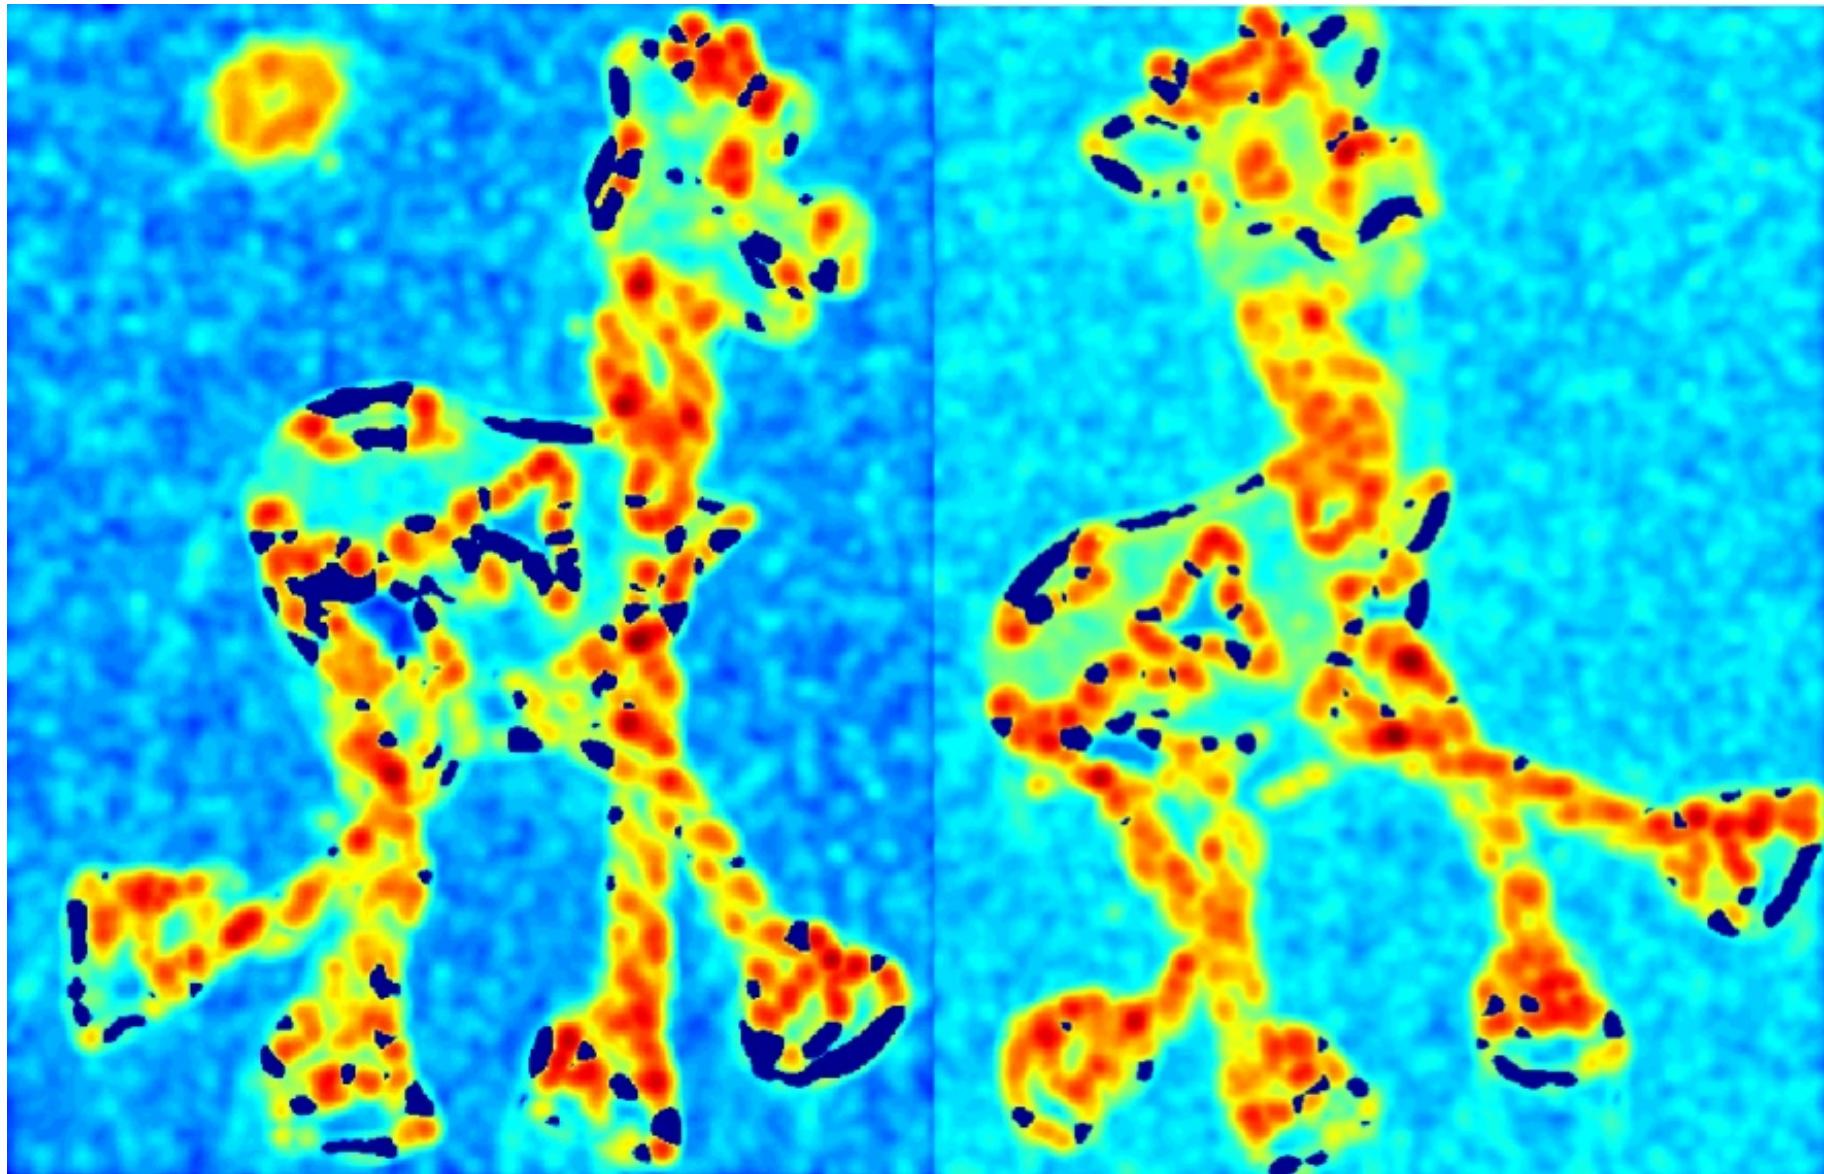
$$f = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2}$$

# Harris corner detector



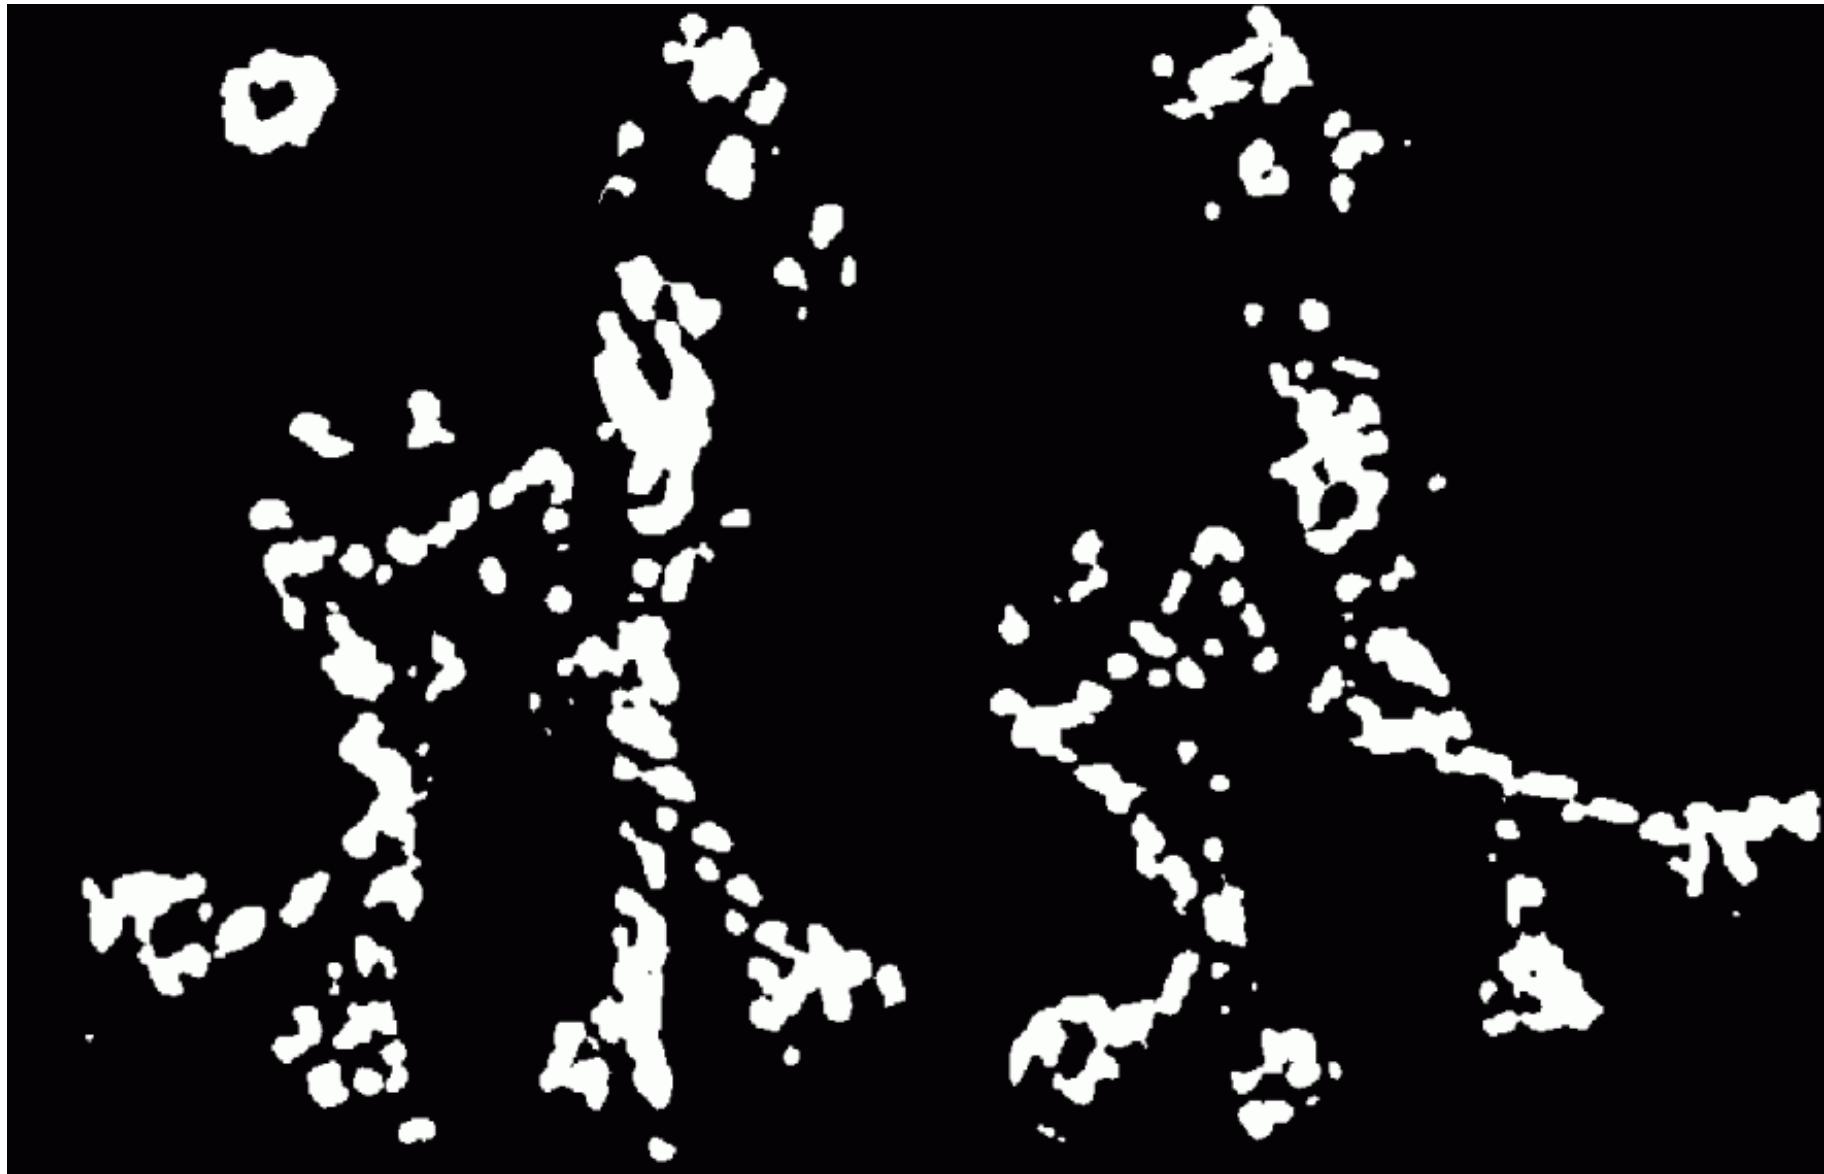
# Harris corner detector

Compute corner response  $f$



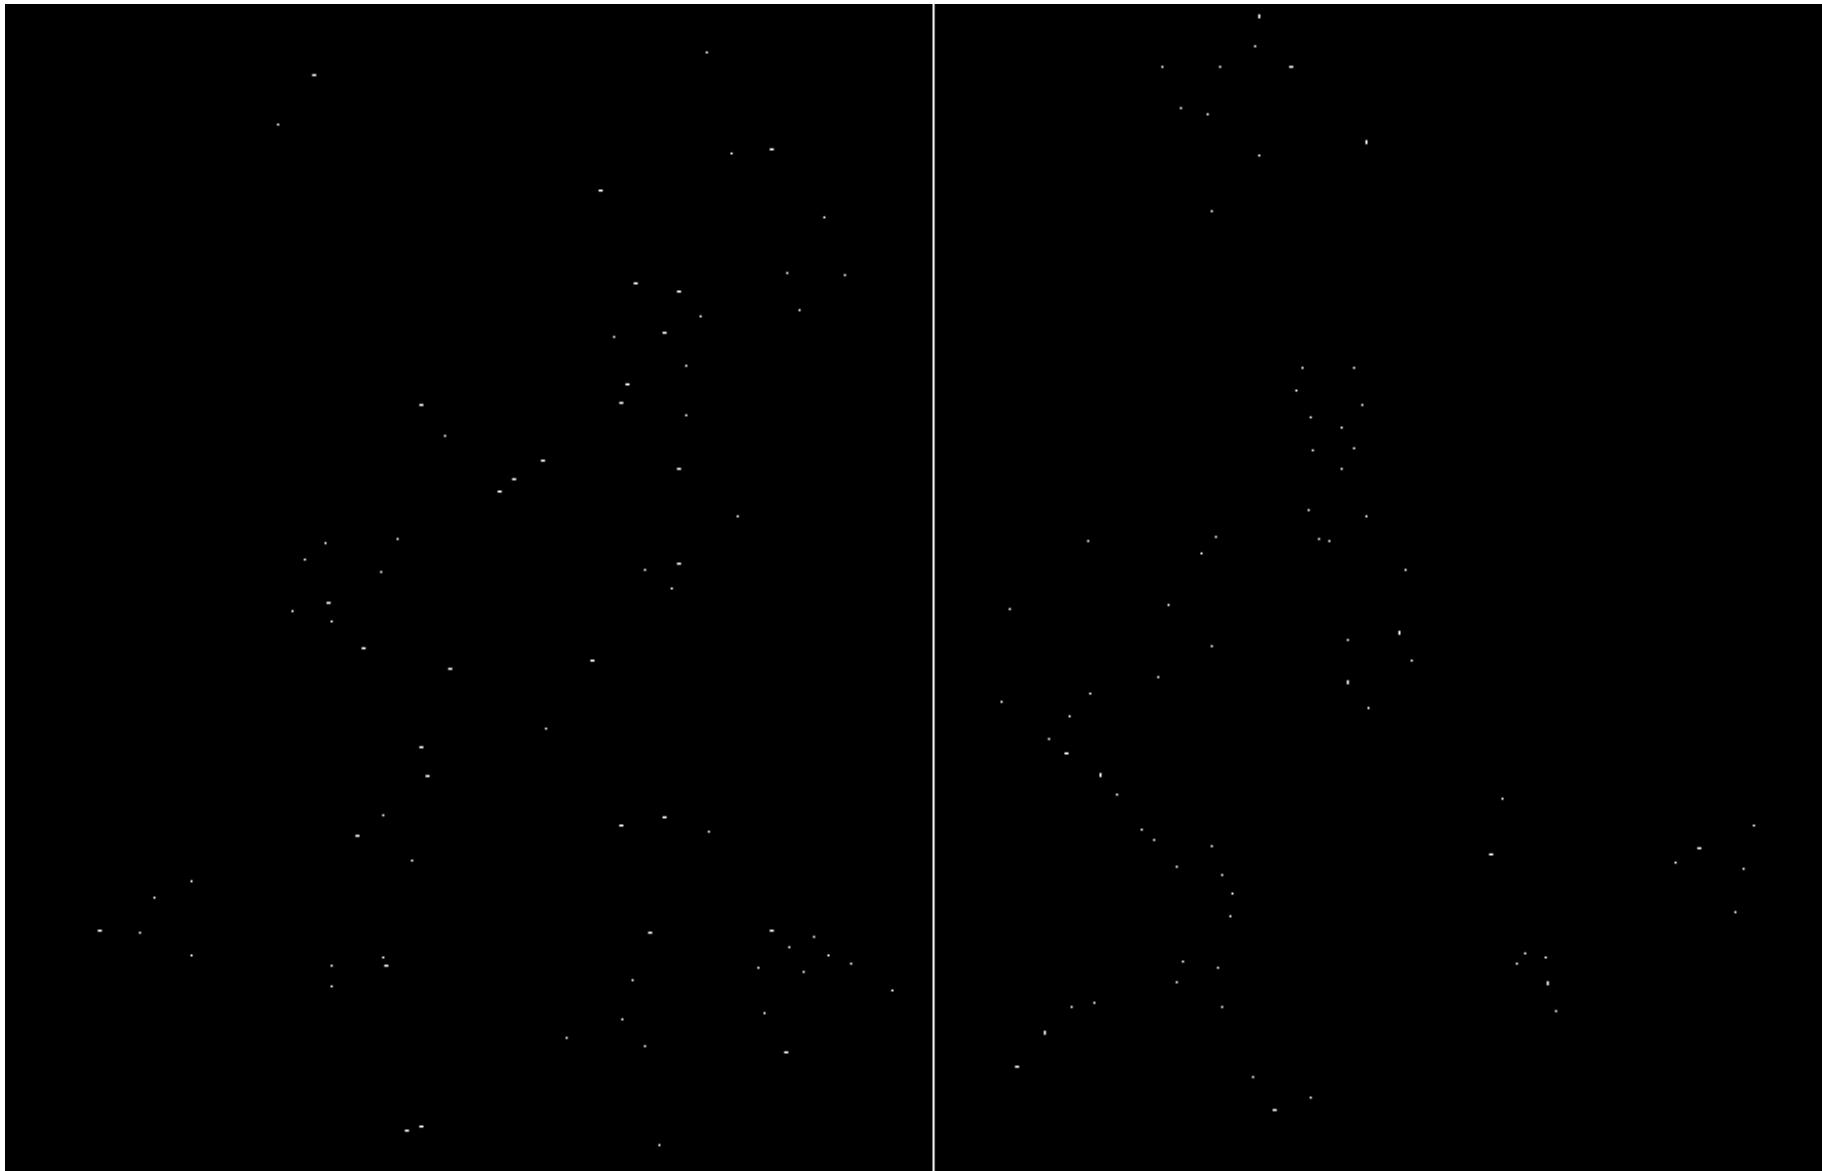
# Harris corner detector

Find points with large corner response:  $f > \text{threshold}$

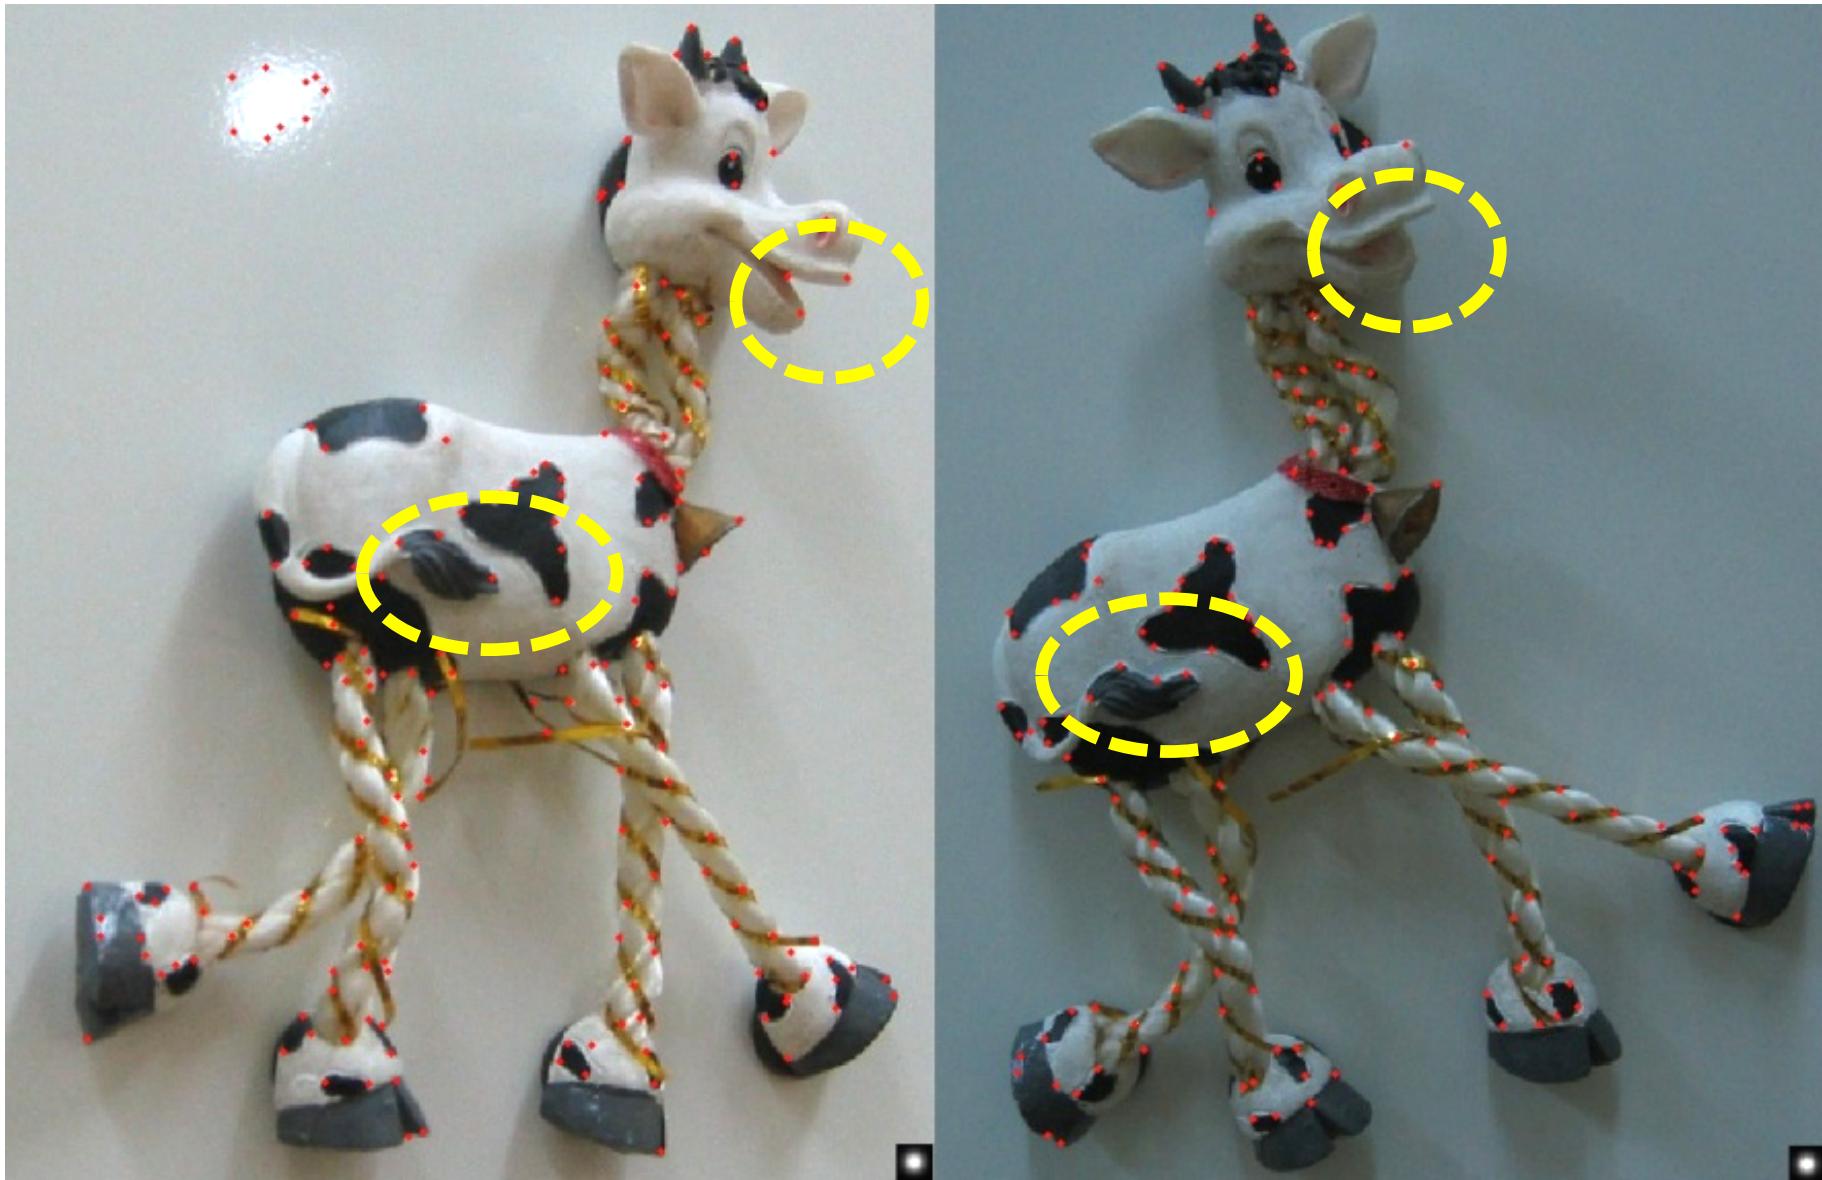


# Harris corner detector

Take only the points of local maxima of  $f$



# Harris corner detector



# Harris corner detector



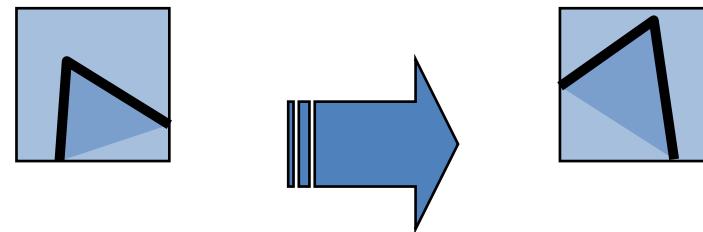
# Harris corner detector



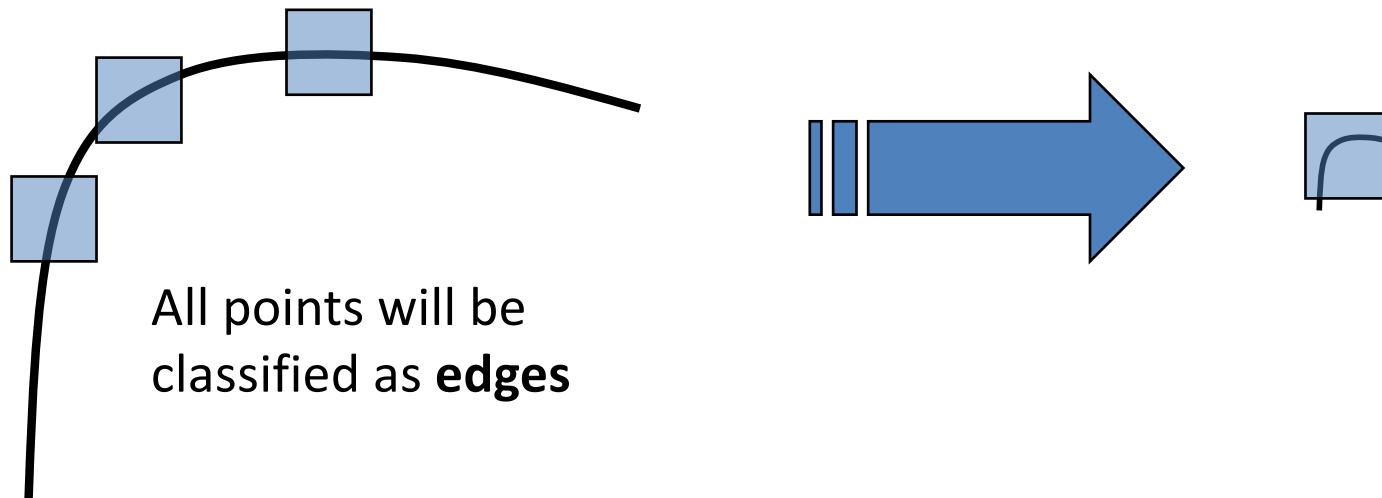
# Harris corner detector

- Properties of the Harris corner detector

- Rotation invariance  
rotates but its shape (i.e.  
eigenvalues) remains the **same**



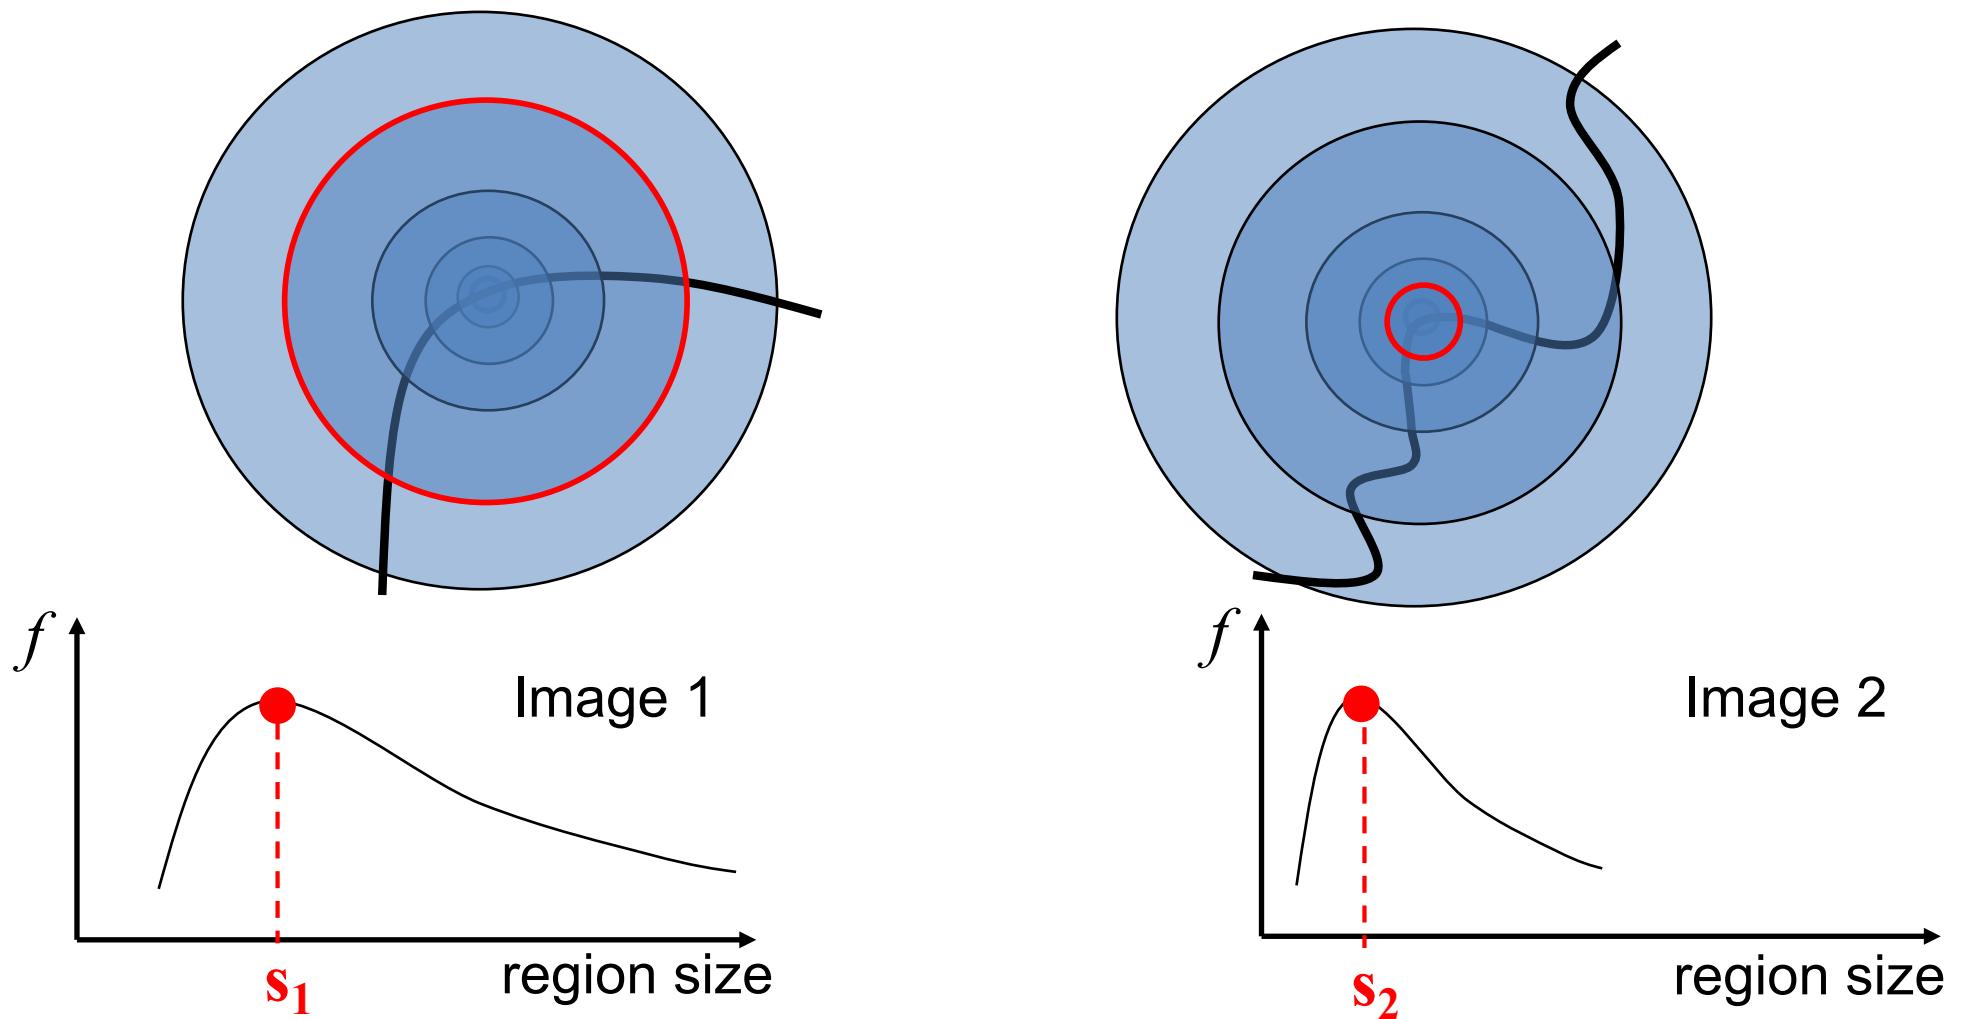
- Not invariant to image scale



- How can we detect **scale invariant** interest points?

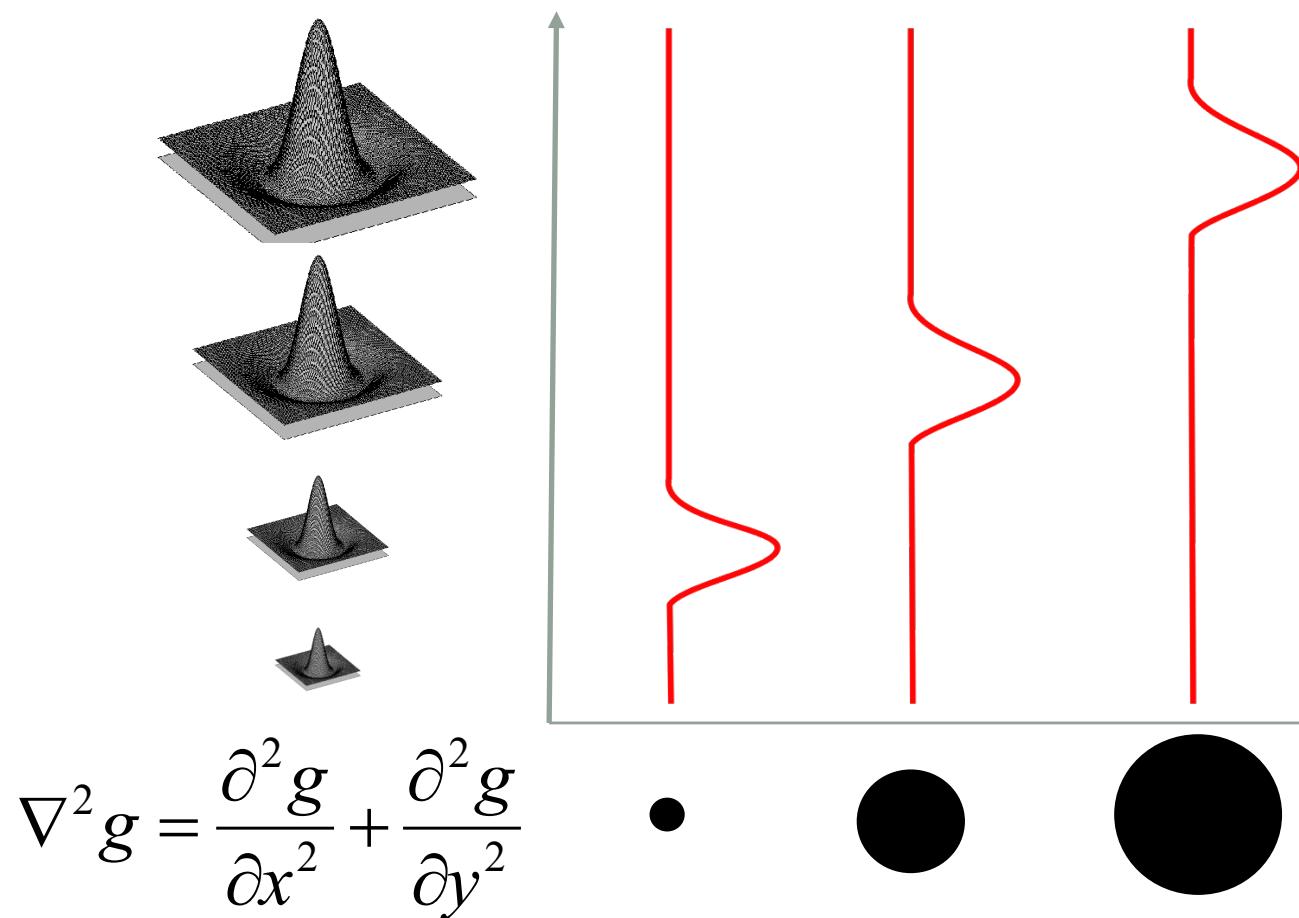
# Automatic scale selection

**Intuition** - Find scale that gives local maxima of some signature function  $f$  in both position and scale.

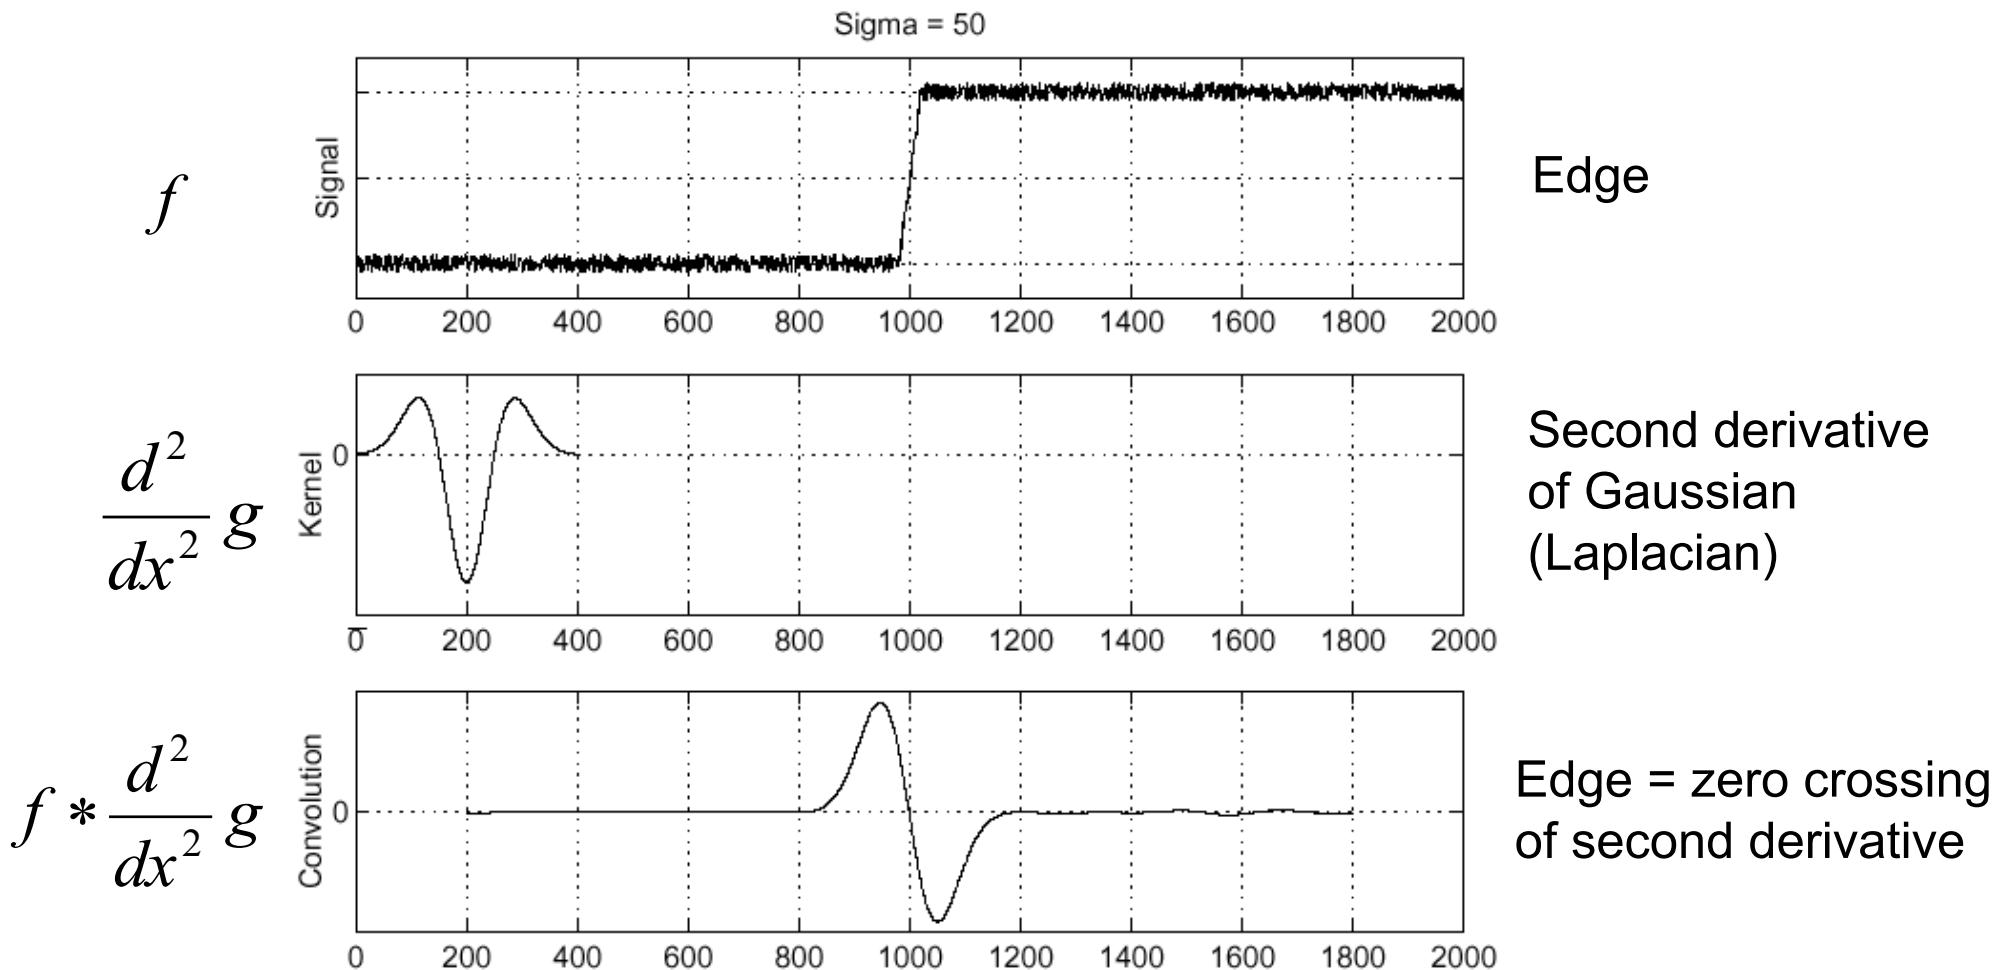


# Scale invariance detection

- Useful signature function
  - Laplacian-of-Gaussian = “blob” detector



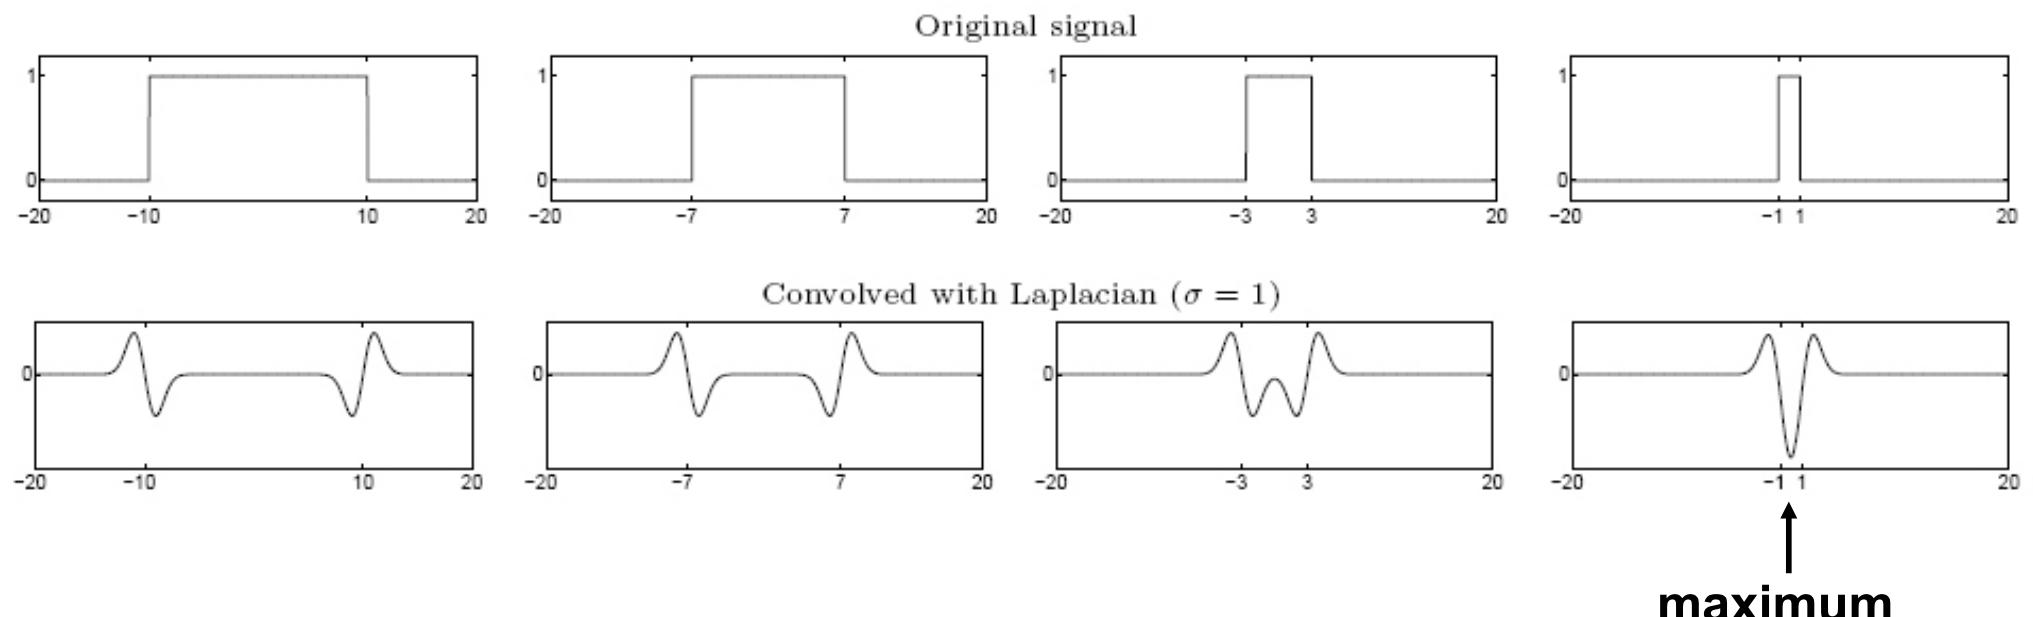
# Scale invariance detection



Source: S. Seitz

# Scale invariance detection

- Edge = ripple
- Blob = superposition of two ripples



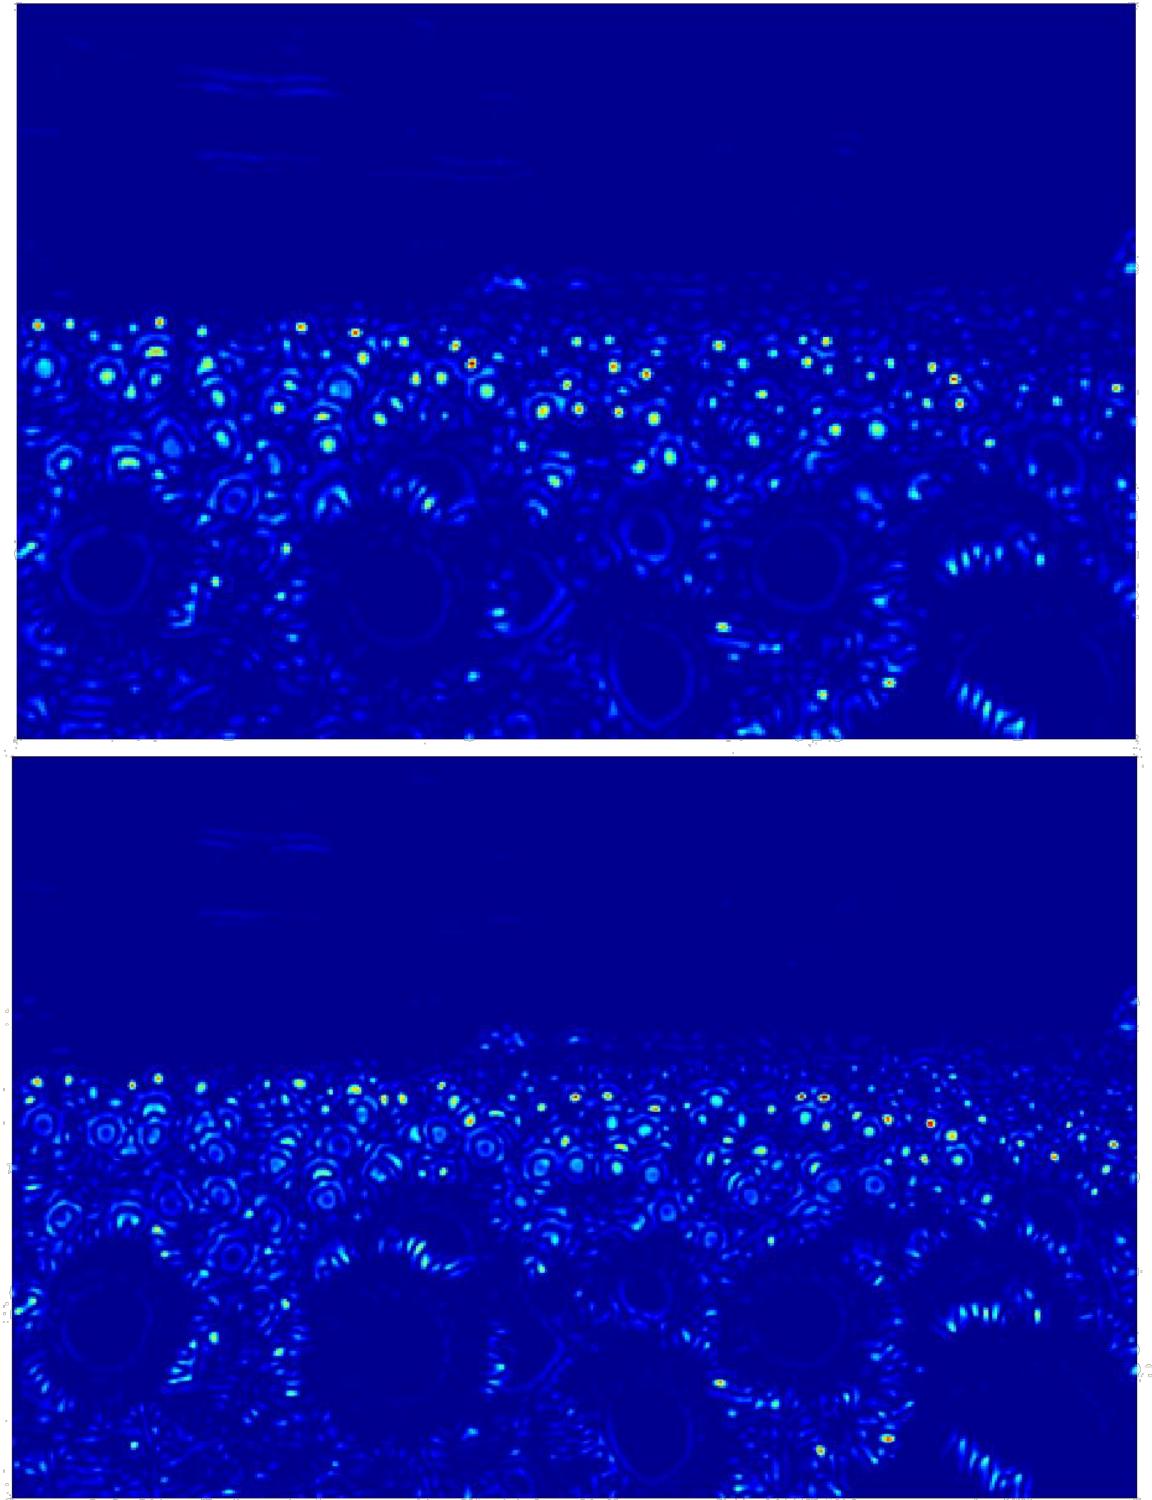
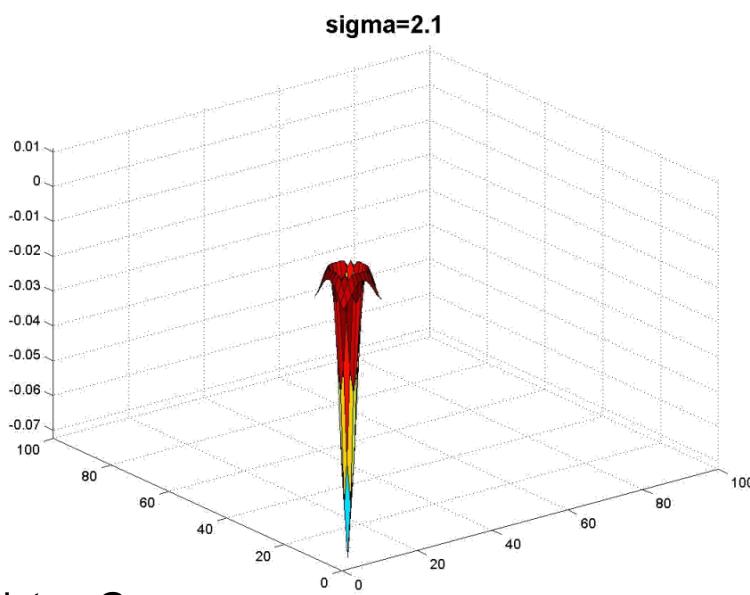
**Spatial selection:** the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

Scaled image  
at  $\frac{3}{4}$  the size

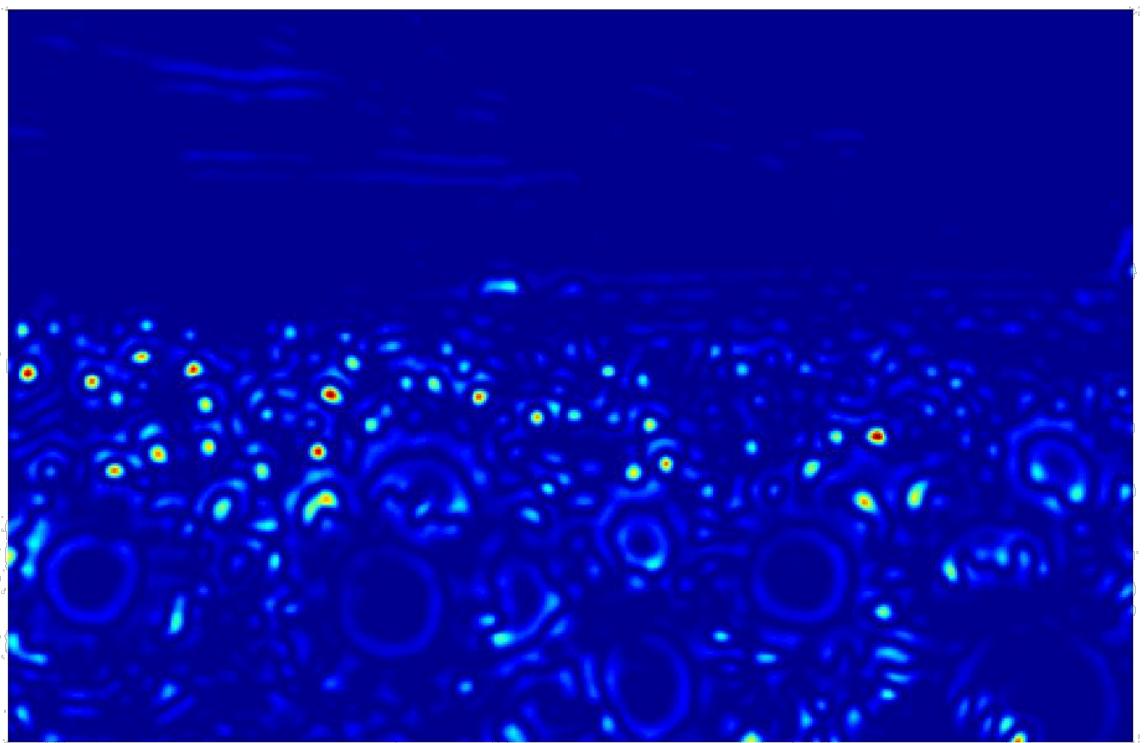
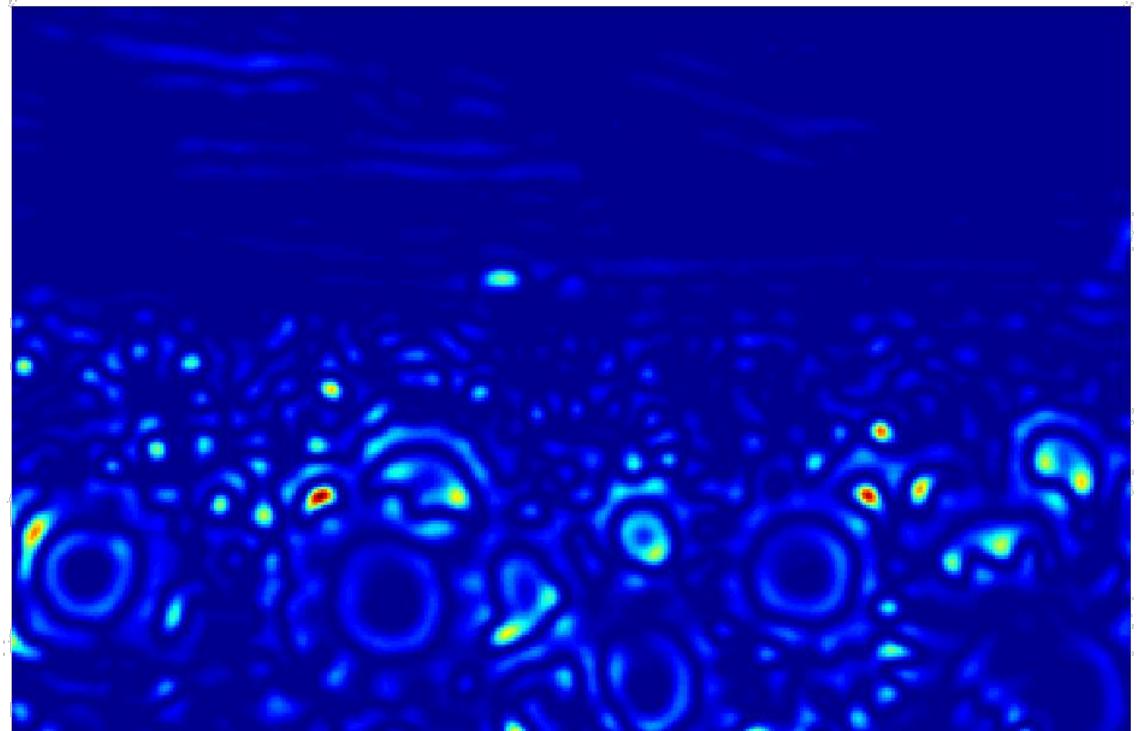
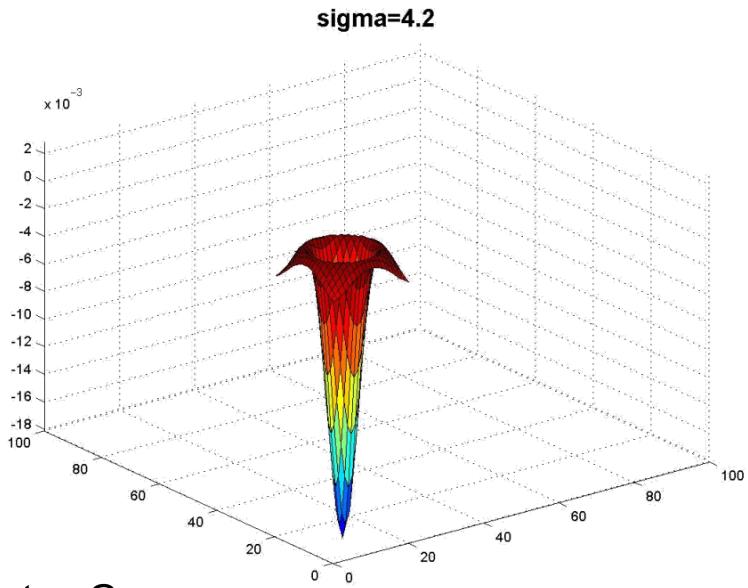


Original image

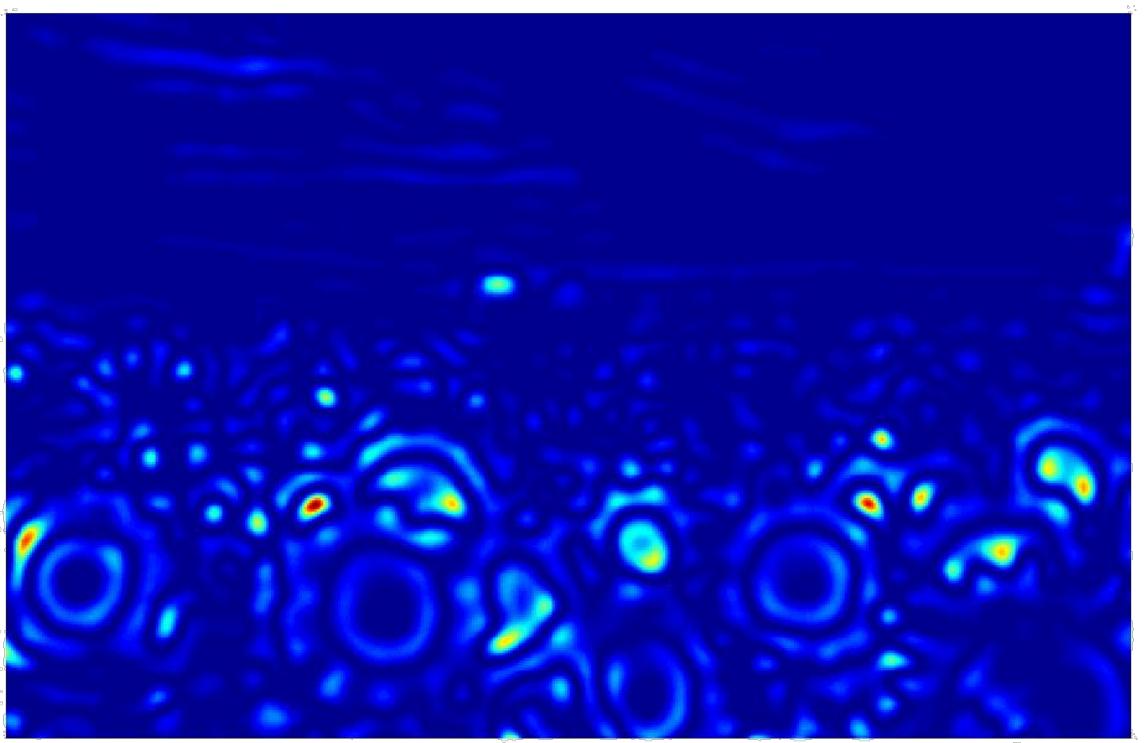
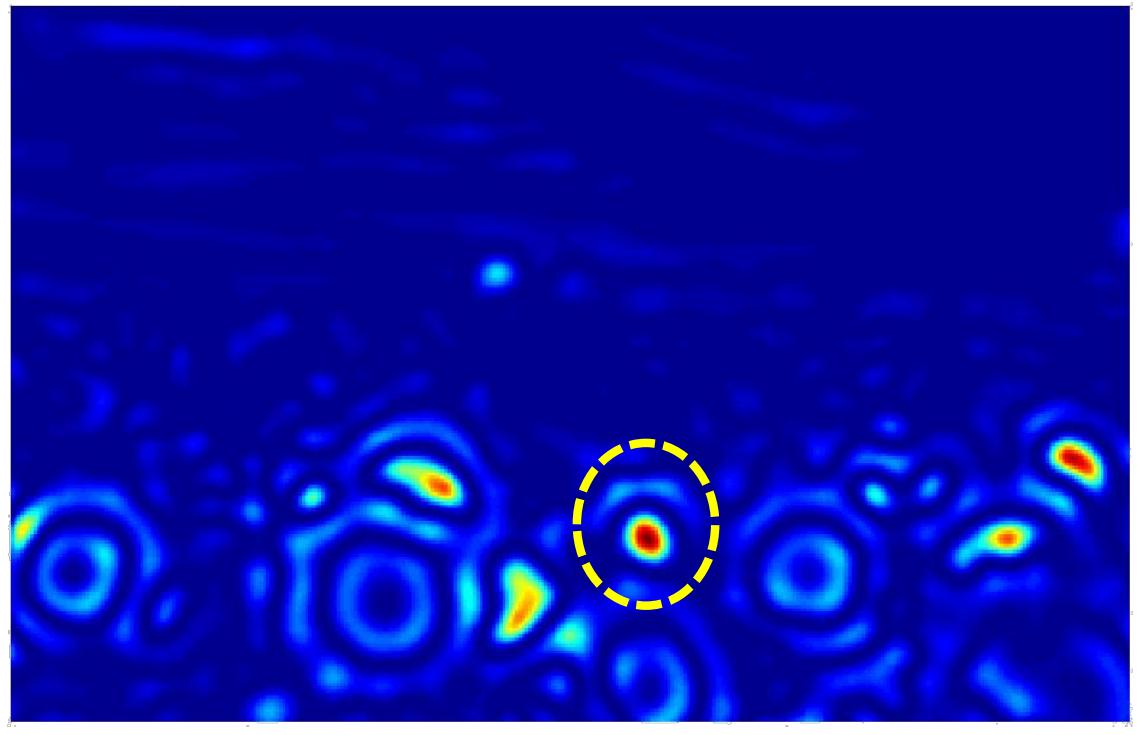
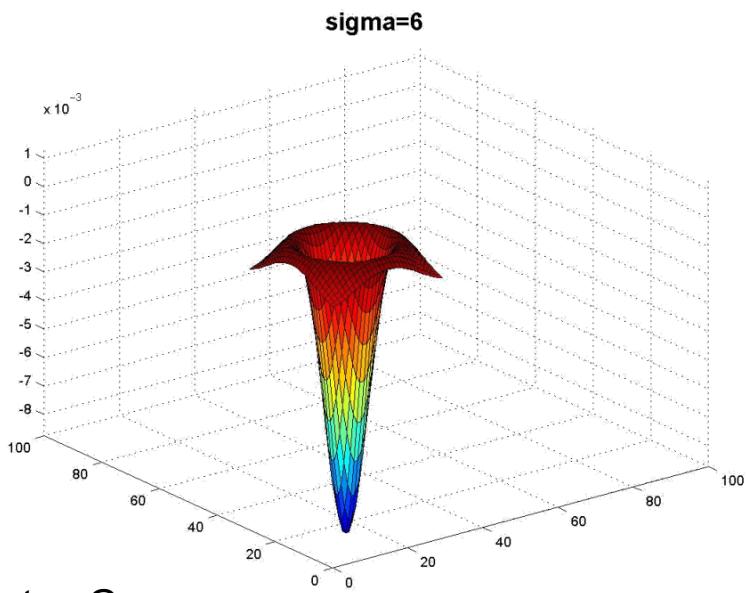




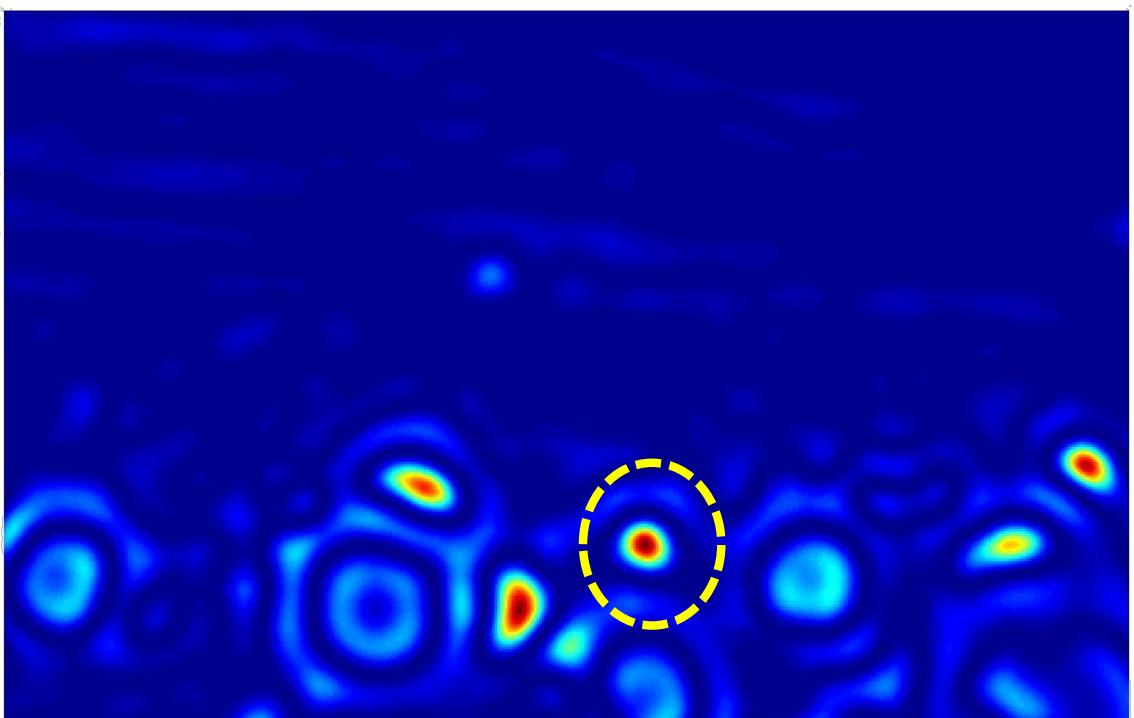
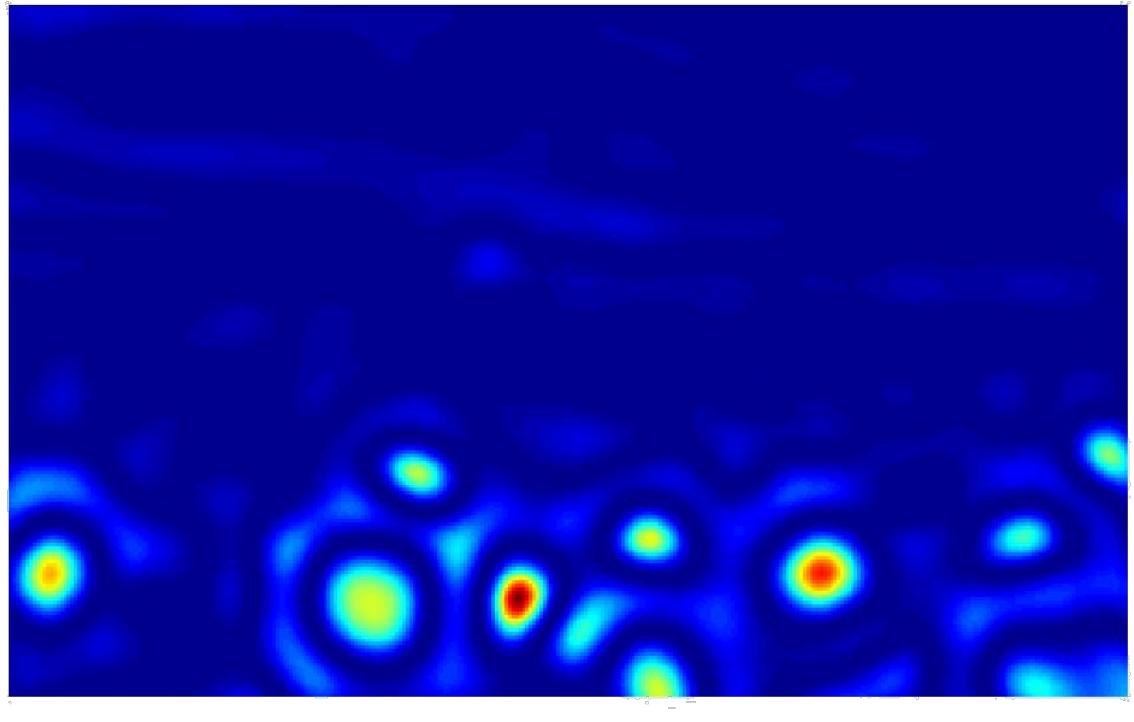
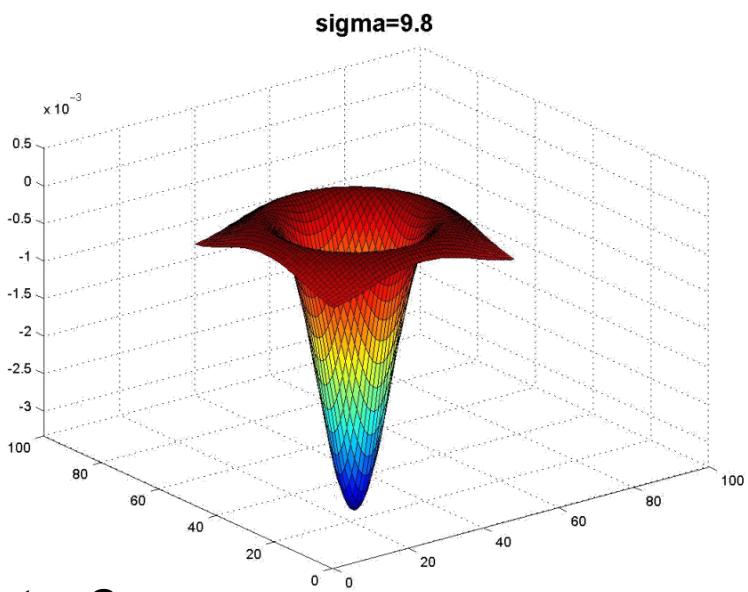
Kristen Grauman

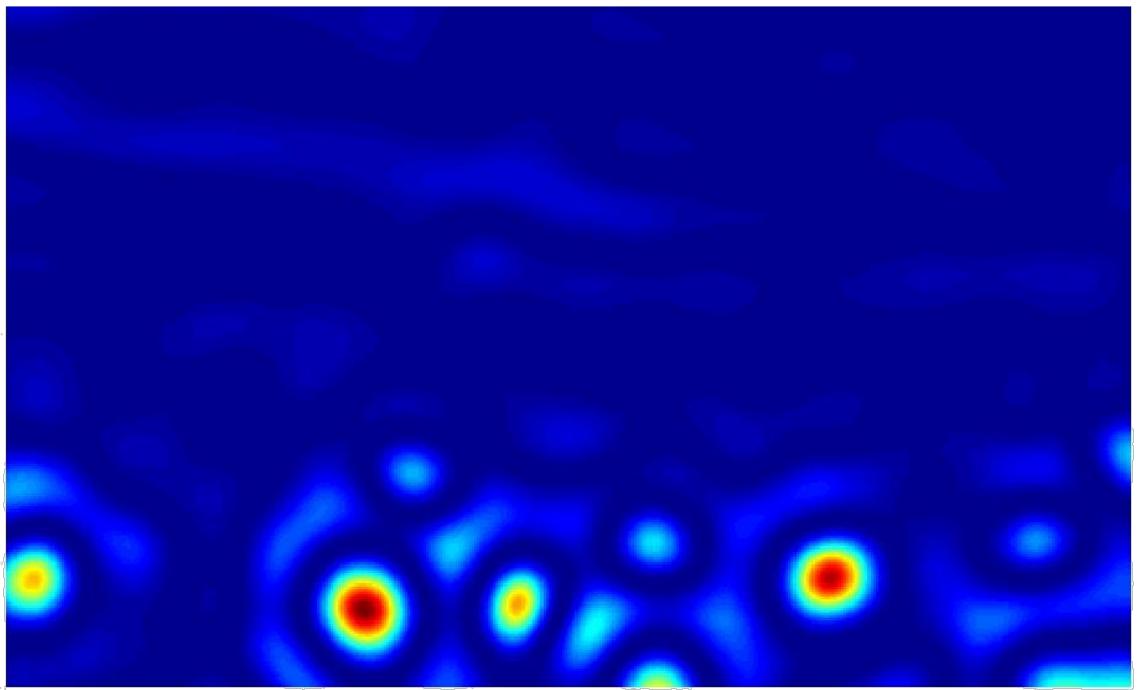
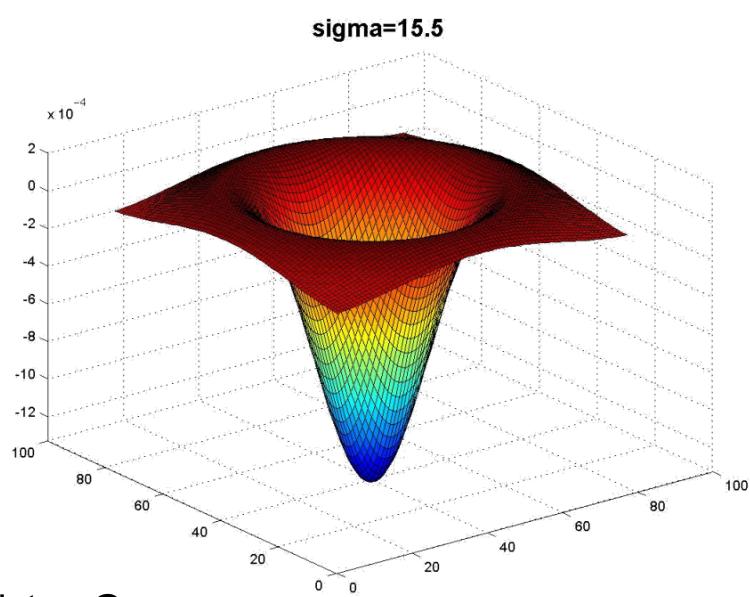
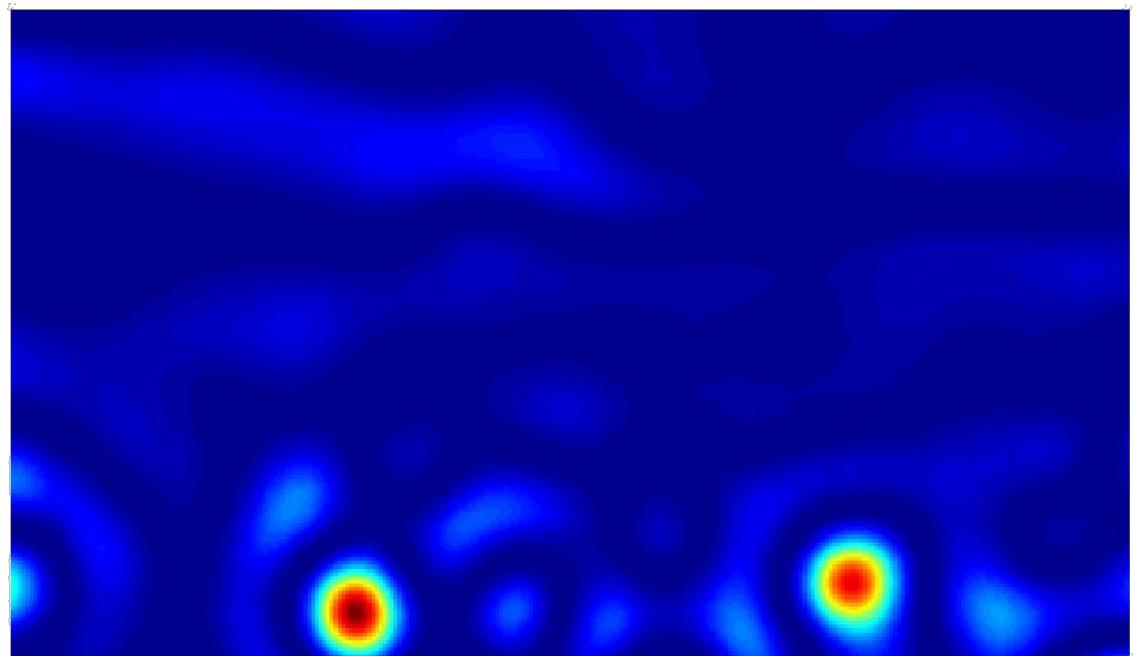


Kristen Grauman

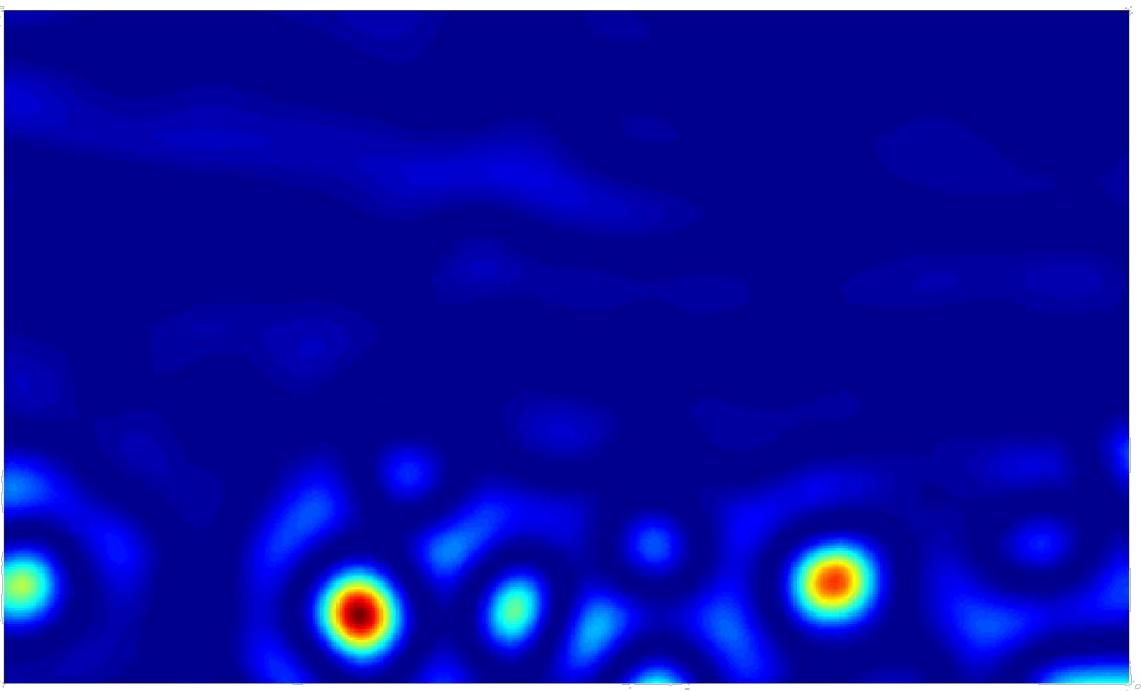
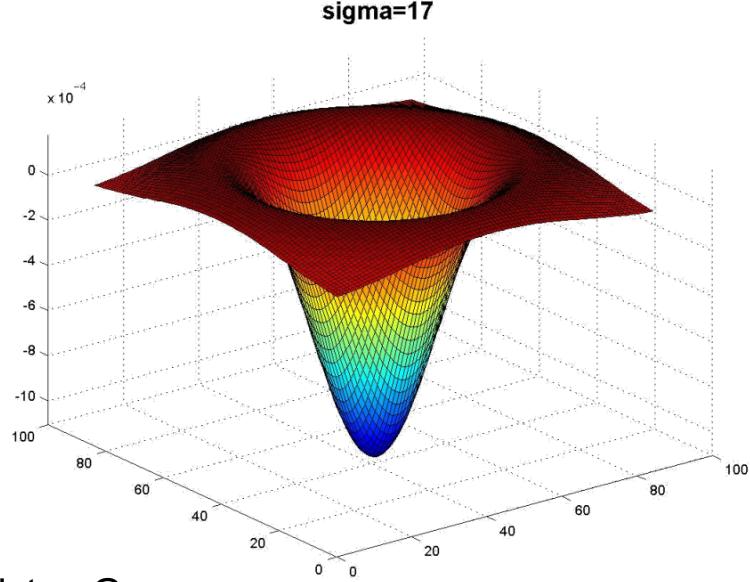
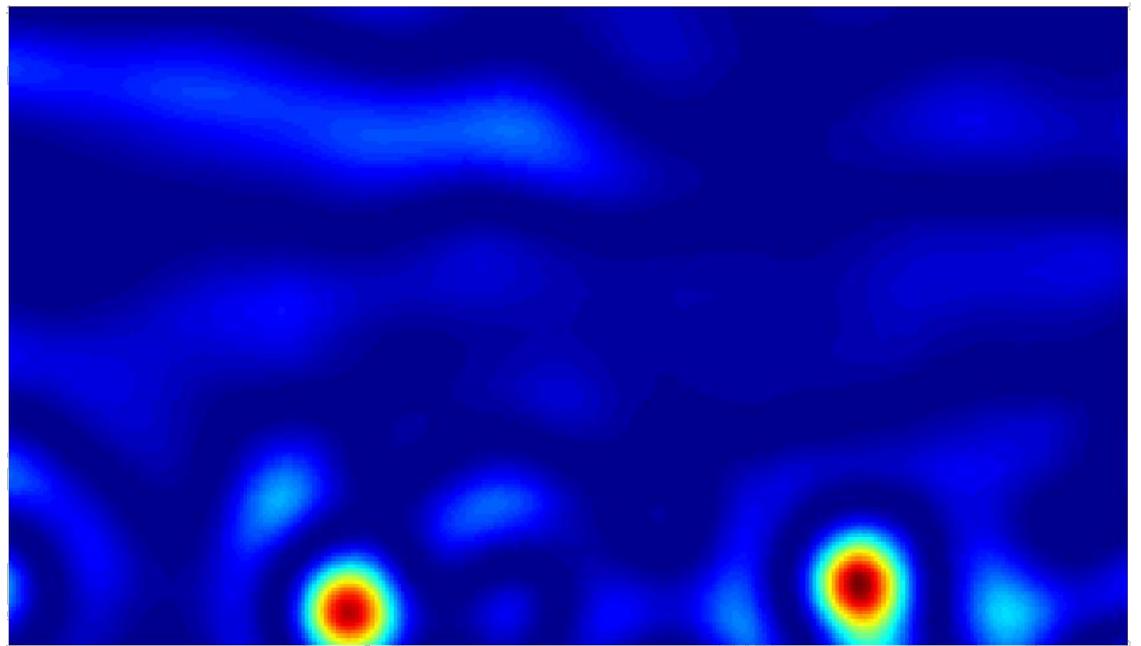


Kristen Grauman





Kristen Grauman



# Scale invariant interest points

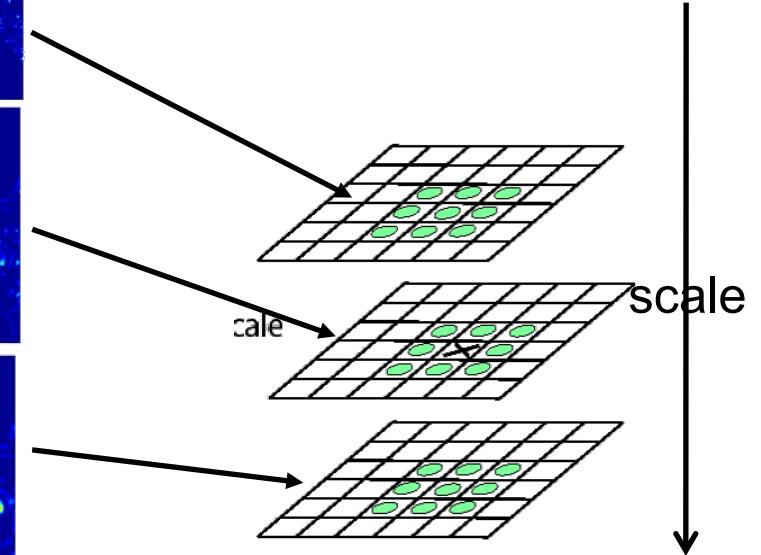
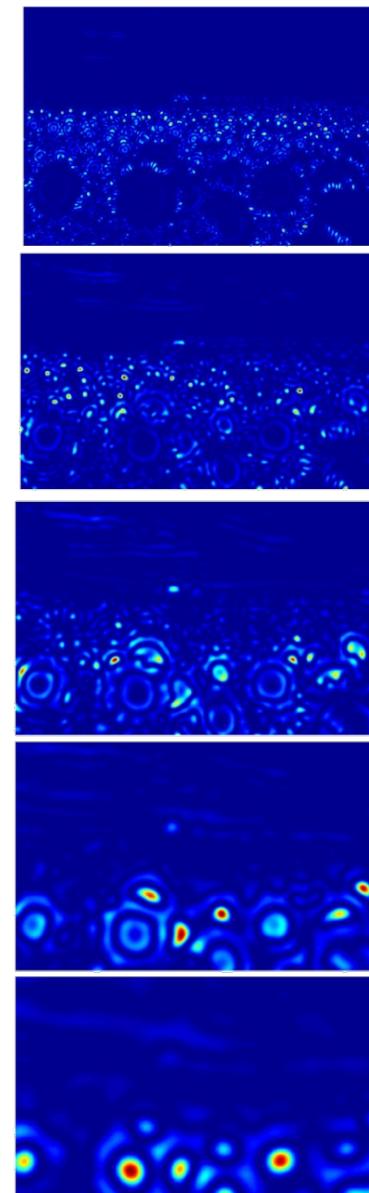
Interest points are local maxima in both position and scale.



$$L_{xx}(\sigma) + L_{yy}(\sigma) \rightarrow \sigma_3$$

Squared filter  
response maps

$$\sigma_5 \\ \sigma_4 \\ \sigma_3 \\ \sigma_2 \\ \sigma_1$$



→ List of  
 $(x, y, \sigma)$

# Scale invariant interest points

- We can approximate the Laplacian with a difference of Gaussians, which is more efficient to implement.

$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

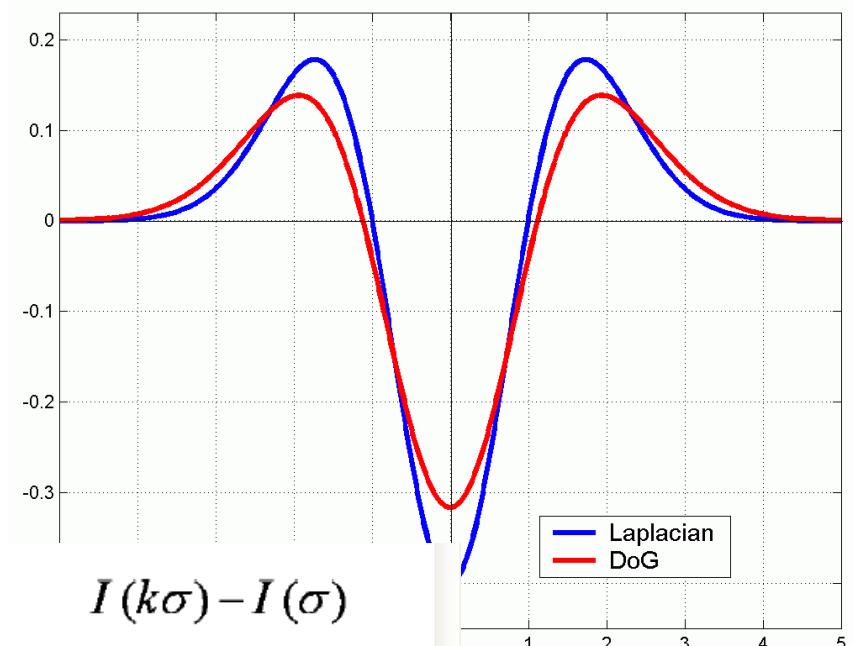
**(Laplacian)**

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

**(Difference of Gaussians)**

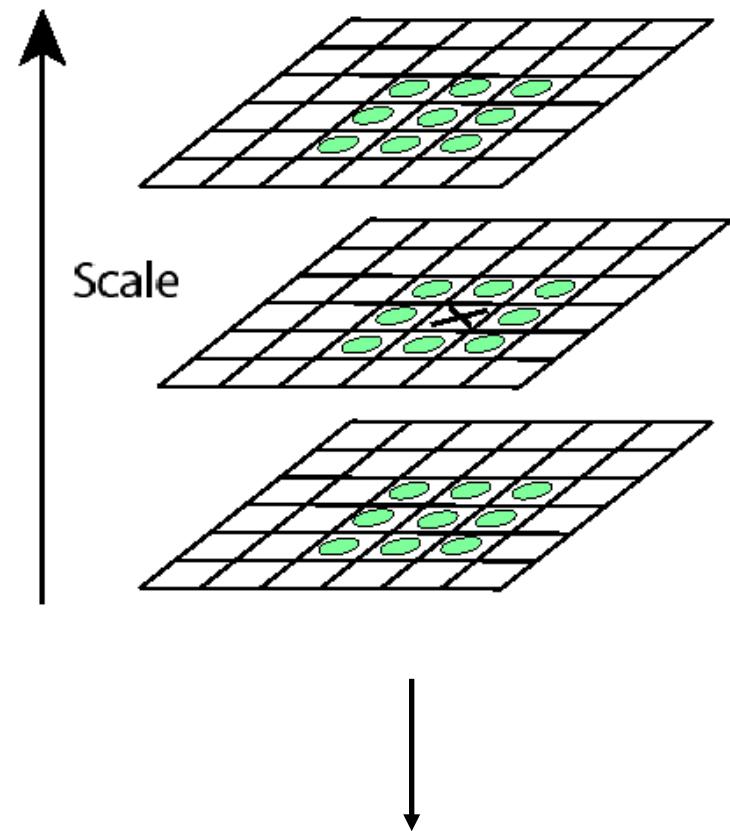


=



# Scale invariance detection

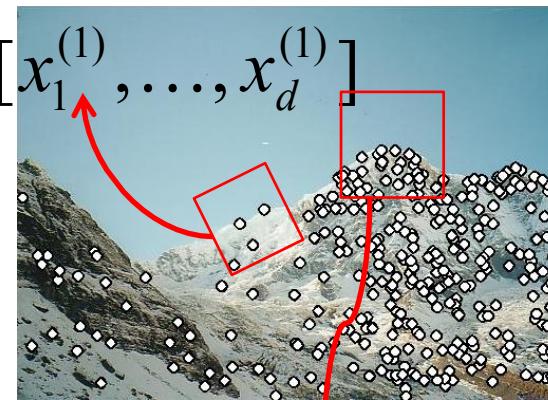
- LoG can be approximated by a Difference of two Gaussians (DoG) at different scales
- Detect maxima of DoG in the scale space volume
- Reject points with low contrast (threshold)
- Reject points that are localized along an edge



Candidate keypoints:  
list of  $(x, y, \sigma)$

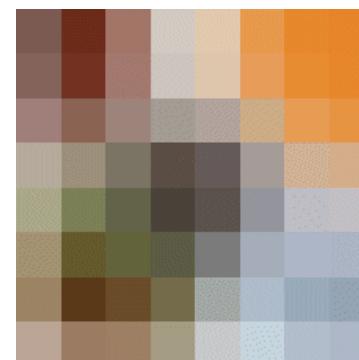
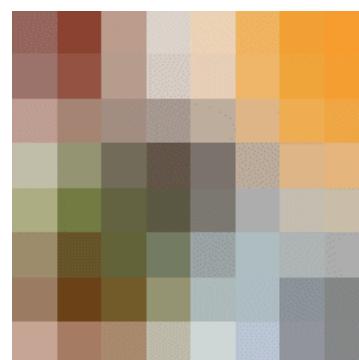
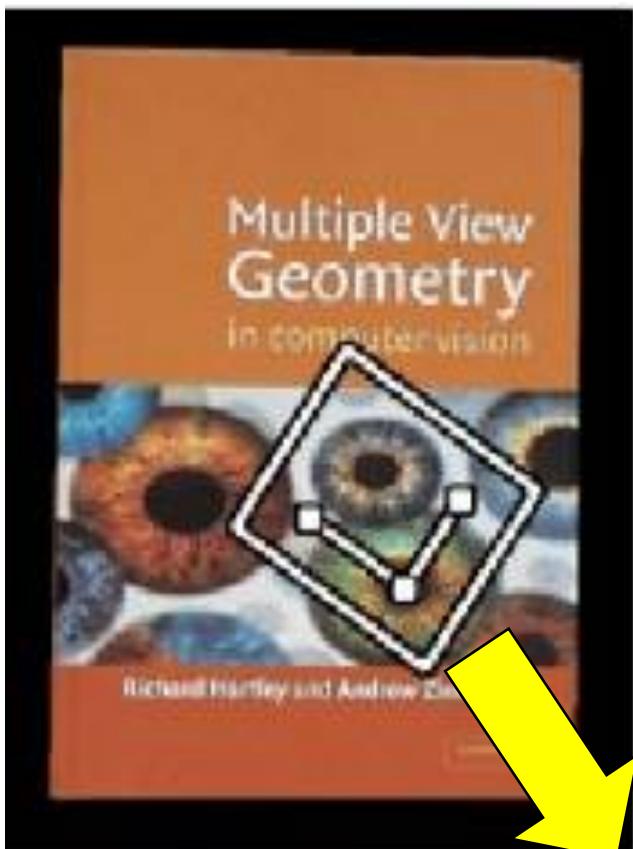
# Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding  $\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$  each interest point.
- 3) Matching: Determine correspondence between descriptors in two views



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

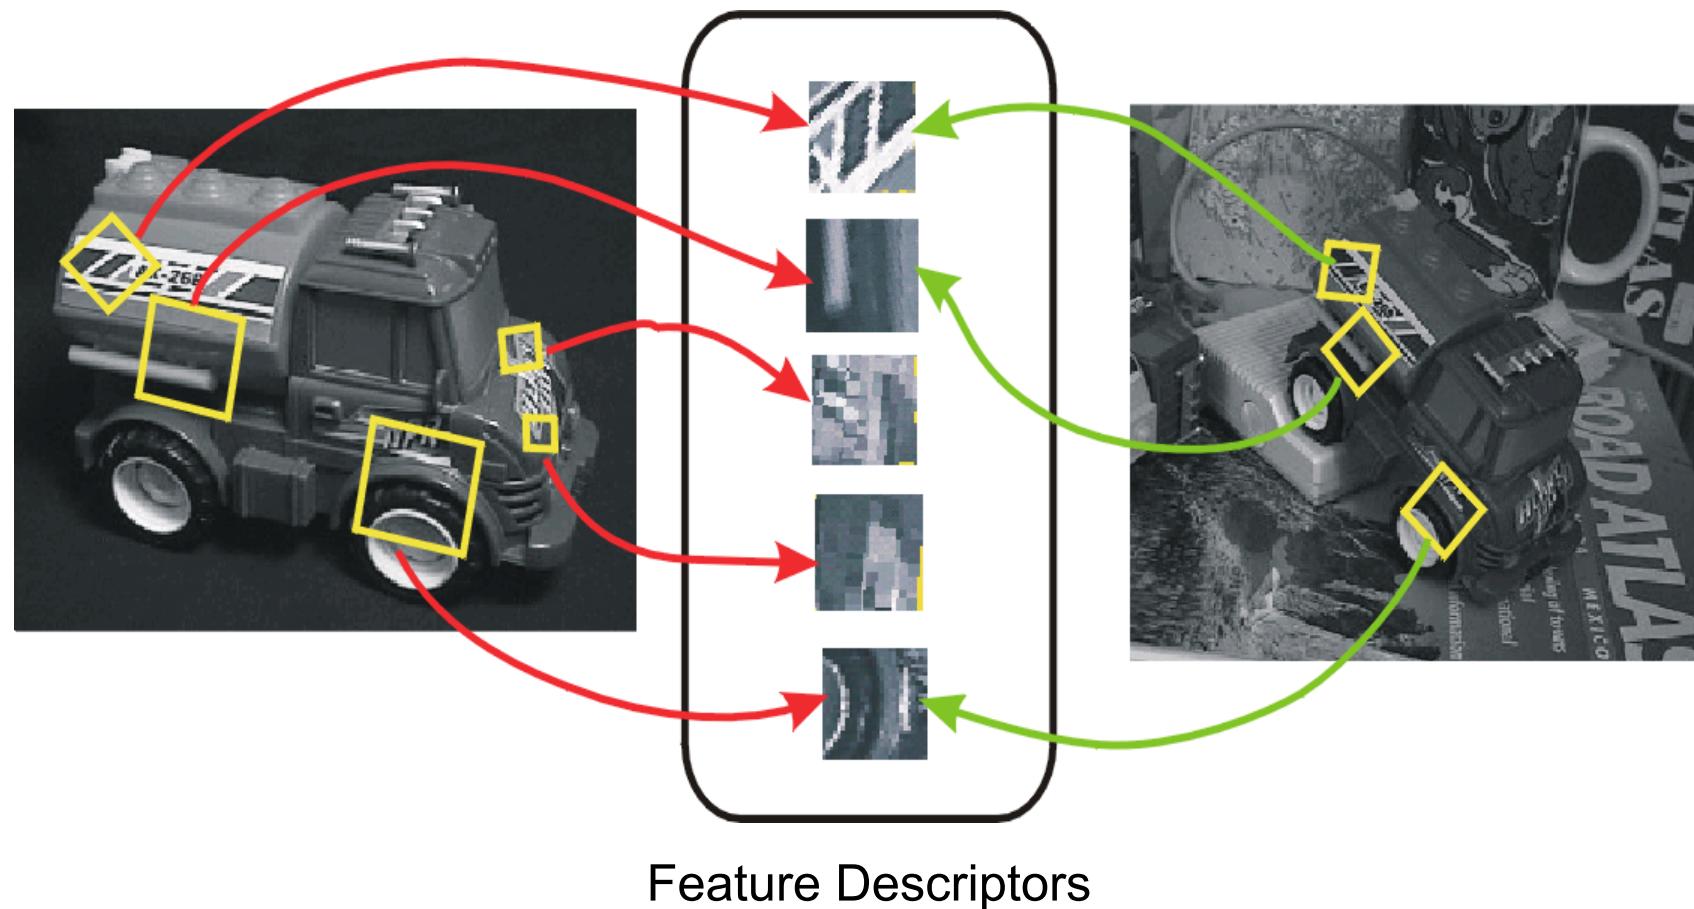
# Geometric transformation



# Invariant local features

Find features that are invariant to transformations

- geometric invariance: translation, rotation, scale
- photometric invariance: brightness, exposure, ...

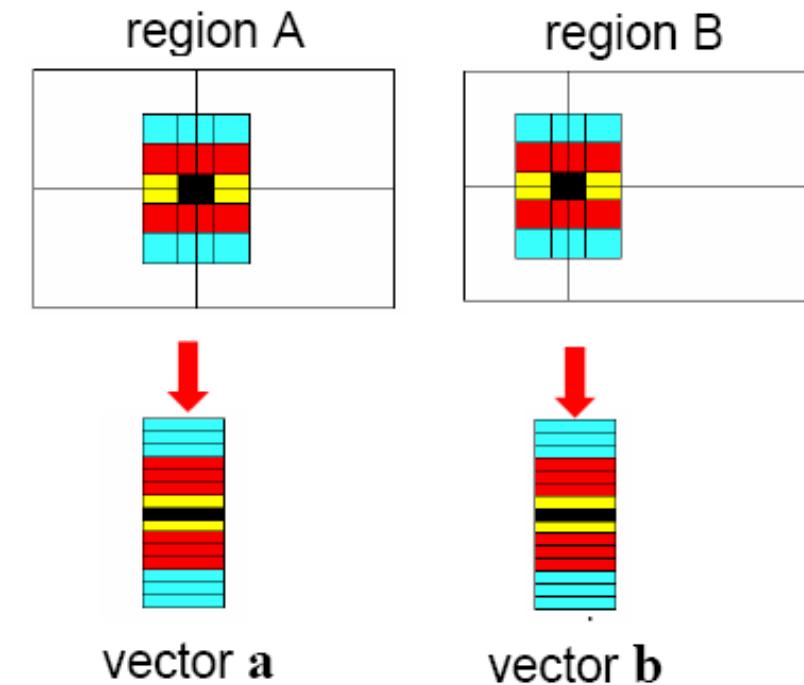
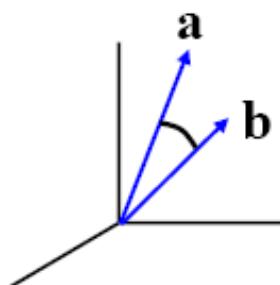


# Local descriptors

- Simplest descriptor: list of intensities within a patch.
- What is this going to be invariant to?

Write regions as vectors

$$A \rightarrow a, B \rightarrow b$$

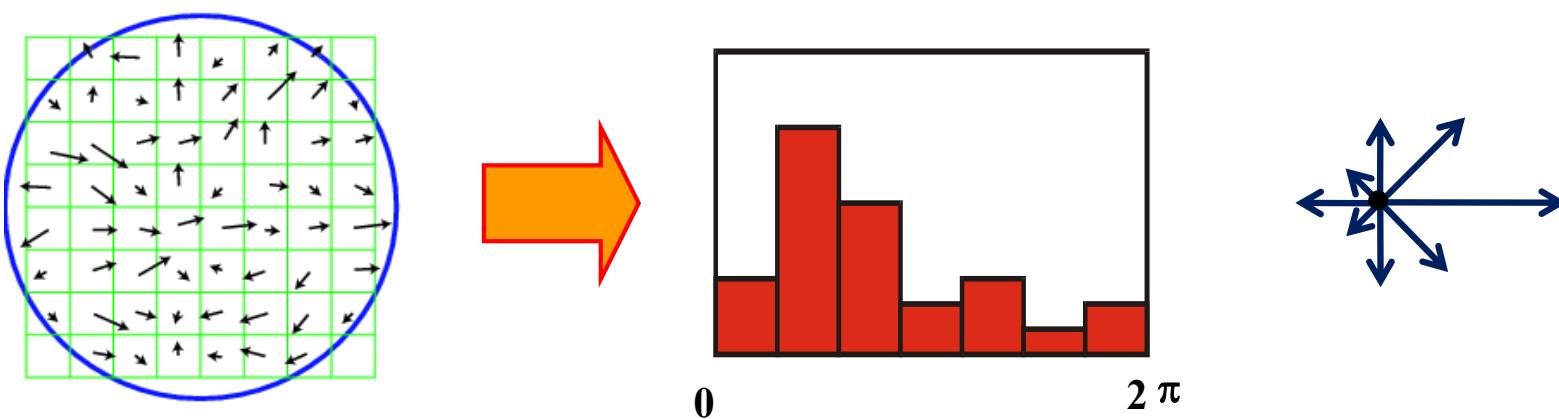


# Local descriptors

- Disadvantage of patches as descriptors:
  - Small shifts can affect matching score a lot



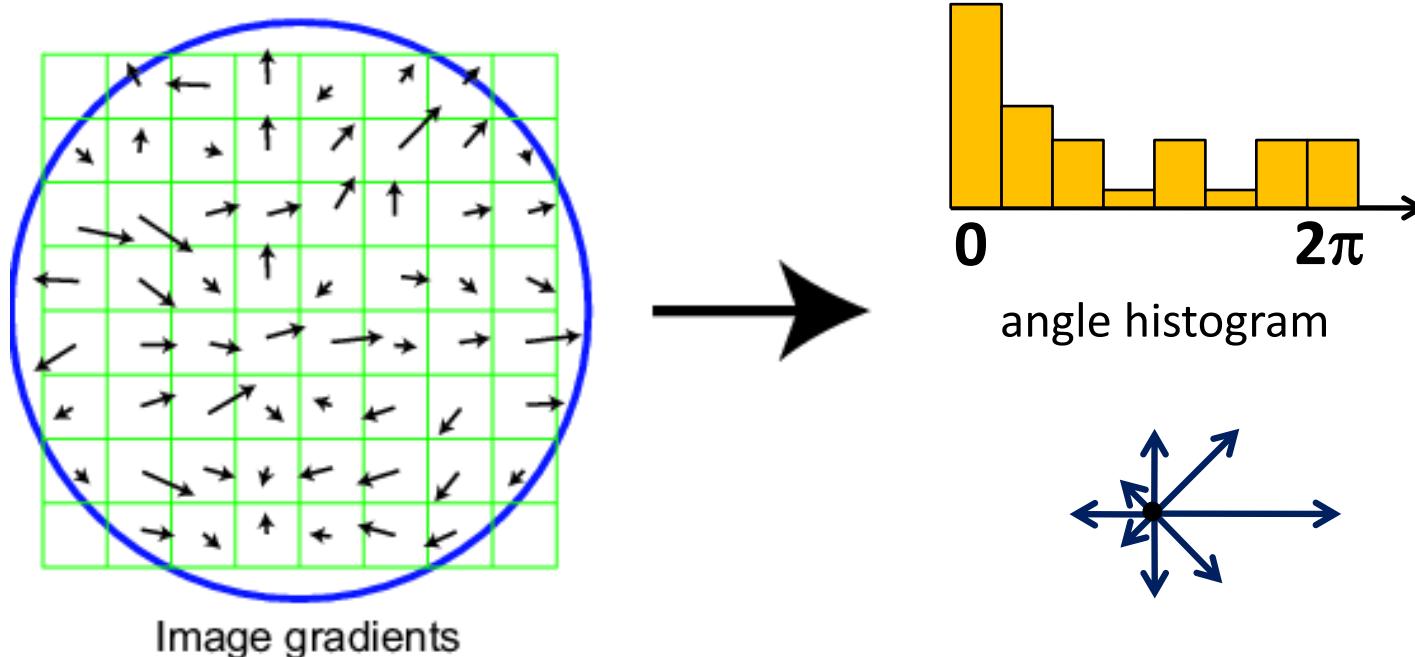
- Solution: histograms



# SIFT descriptor

Basic idea:

- Take 16x16 square window around detected feature
- Compute edge orientation (angle of the gradient -  $90^\circ$ ) for each pixel
- Throw out weak edges (threshold gradient magnitude)
- Create histogram of surviving edge orientations

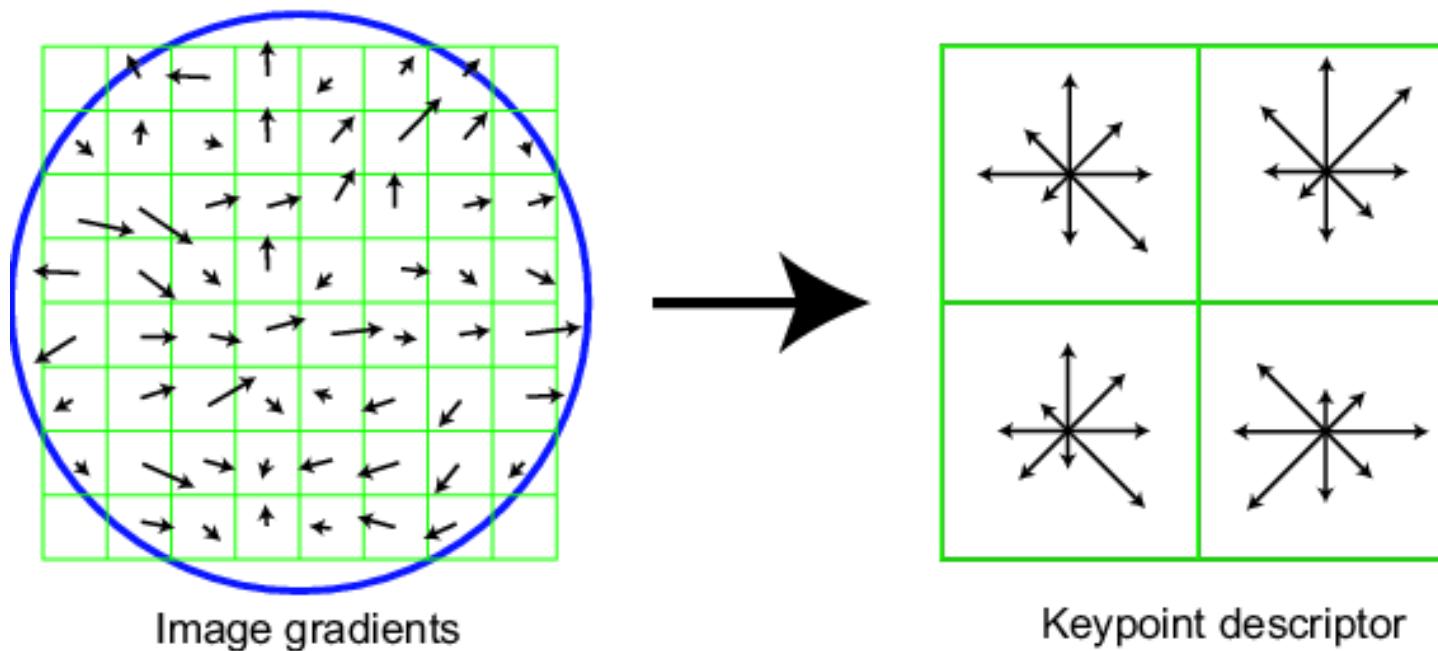


Distinctive image features from scale-invariant keypoints. David G. Lowe. *IJCV* 60 (2), pp. 91-110, 2004.

# SIFT descriptor

## Full version

- Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
- Compute an orientation histogram for each cell
- 16 cells \* 8 orientations = 128 dimensional descriptor



# SIFT summary

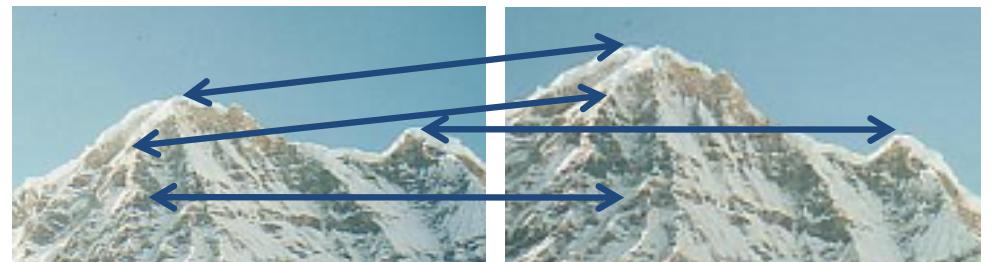
- Sparse, distinctive point features
- Sometimes unclear what features correspond to in image
- **Translation** independent by using **local histogram**
- **Rotation** independent by **orientation adjustment**
- **Scale** independent by **extremal scale estimation**
- **Illumination** independent by **descriptor normalisation**
- Widely used
- Real-time implementation possible

# Local descriptors

- Histogram-based descriptors
  - Based on the histogram of oriented gradient
  - SIFT, SURF, GLOH and HOG
- Compact descriptors
  - Based on binary strings obtained comparing pairs of image intensities
  - BRIEF, ORB, BRISK and FREAK

# Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views



# Feature matching

Given a feature in  $I_1$ , how to find the best match in  $I_2$ ?

1. Define distance function that compares two descriptors
2. Test all the features in  $I_2$ , find the one with min distance



$I_1$

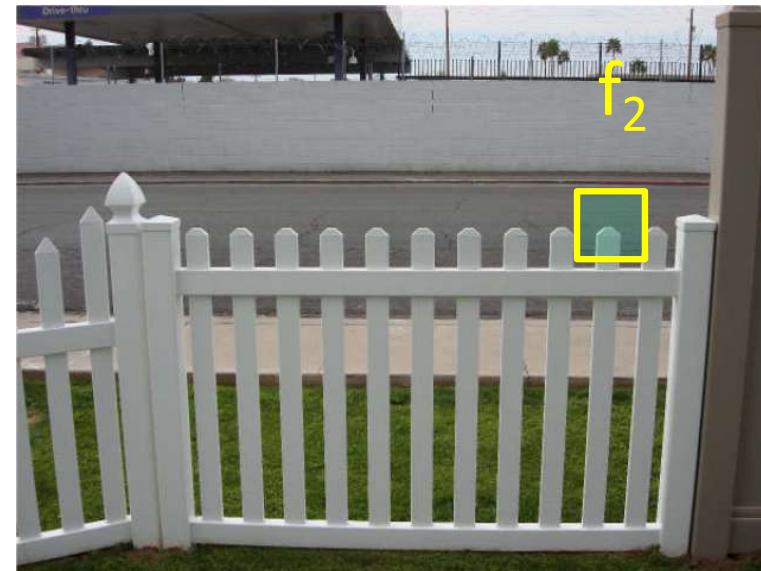
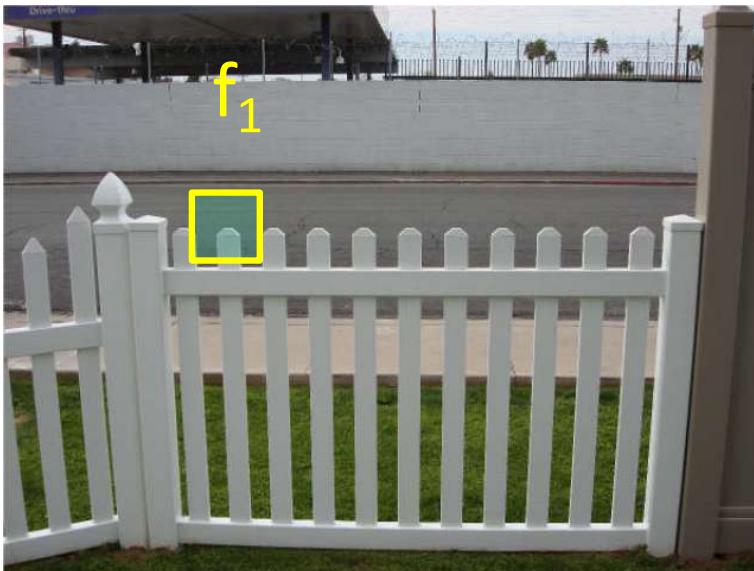


$I_2$

# Feature matching

How to define the difference between two features  $f_1, f_2$ ?

- Simple approach is  $\text{SSD}(f_1, f_2)$ 
  - sum of square differences between entries of the two descriptors
  - can give good scores to very ambiguous (bad) matches

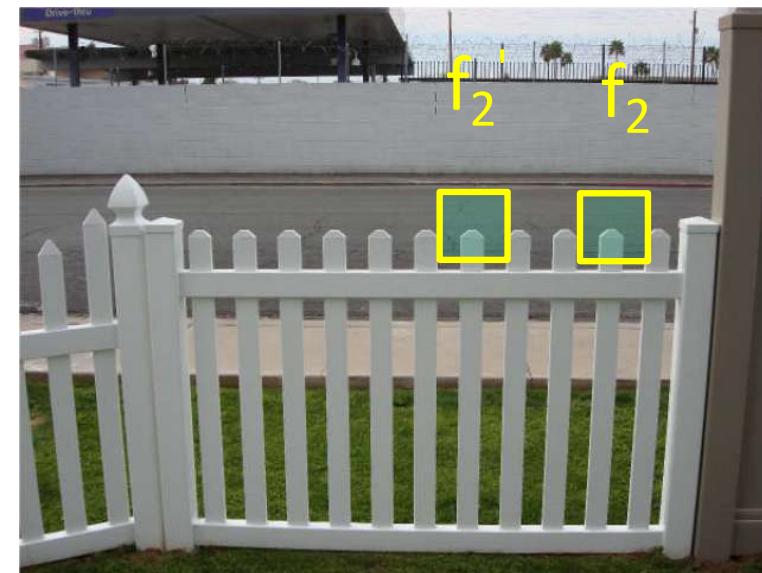
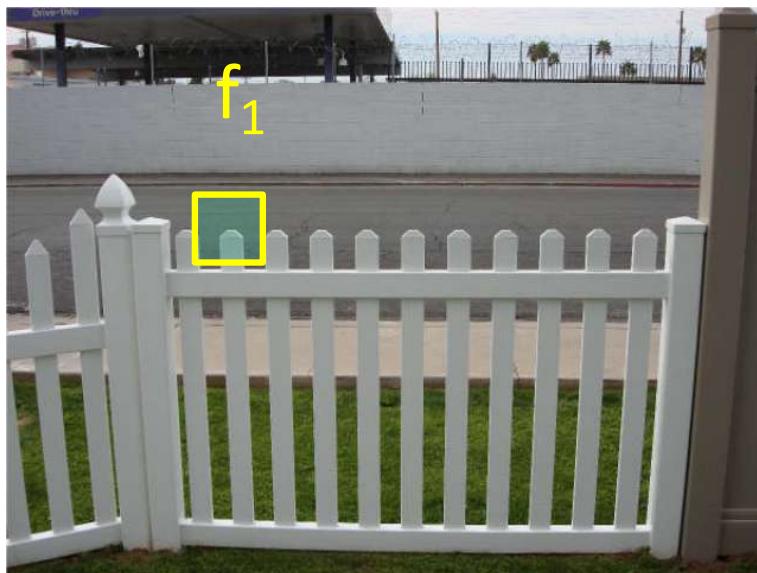


# Feature matching

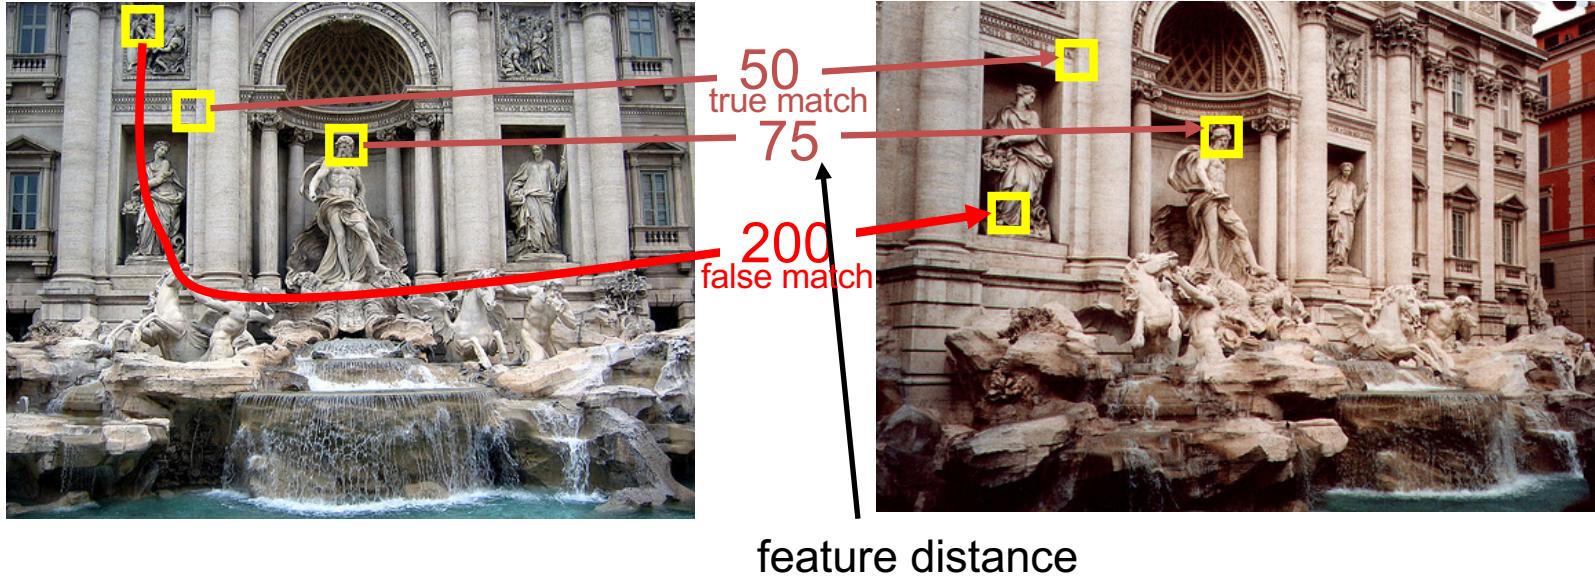
How to define the difference between two features  $f_1, f_2$ ?

- Ratio distance =  $\text{SSD}(f_1, f_2) / \text{SSD}(f_1, f_2')$

- $f_2$  is best SSD match to  $f_1$  in  $I_2$
- $f_2'$  is 2<sup>nd</sup> best SSD match to  $f_1$  in  $I_2$
- gives large values ( $\sim 1$ ) for ambiguous matches



# Feature matching



- Eliminate **bad matches**: throw out features with distance > threshold

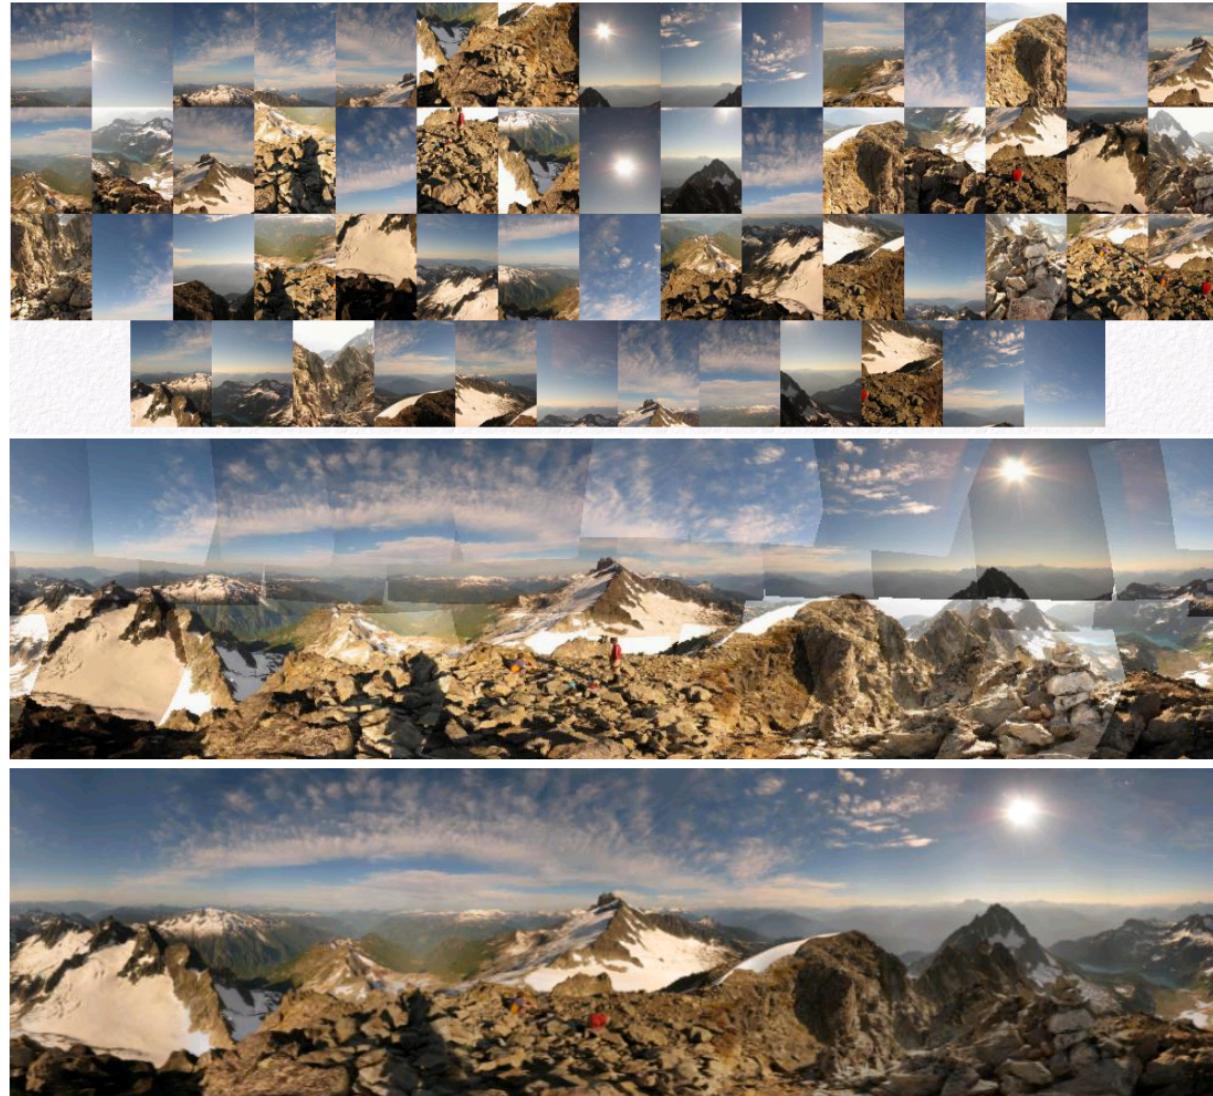
# Applications

Features are used for:

- Image alignment (e.g., mosaics)
- 3D reconstruction
- Motion tracking
- Object recognition
- Indexing and database retrieval
- Robot navigation
- ... other

# Applications

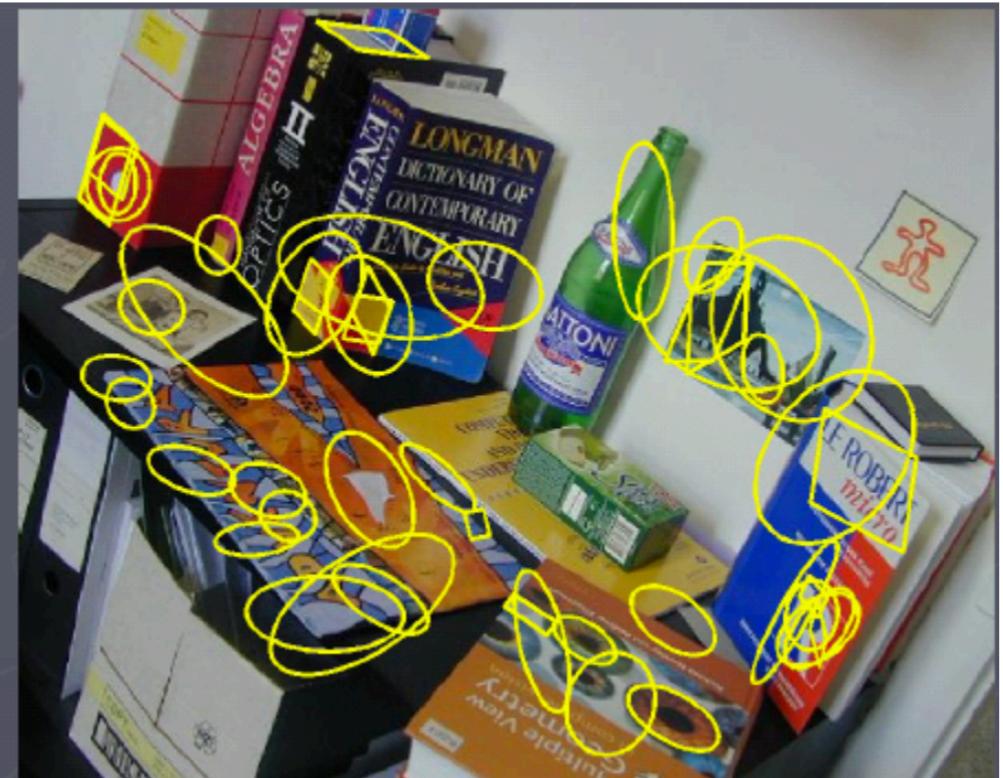
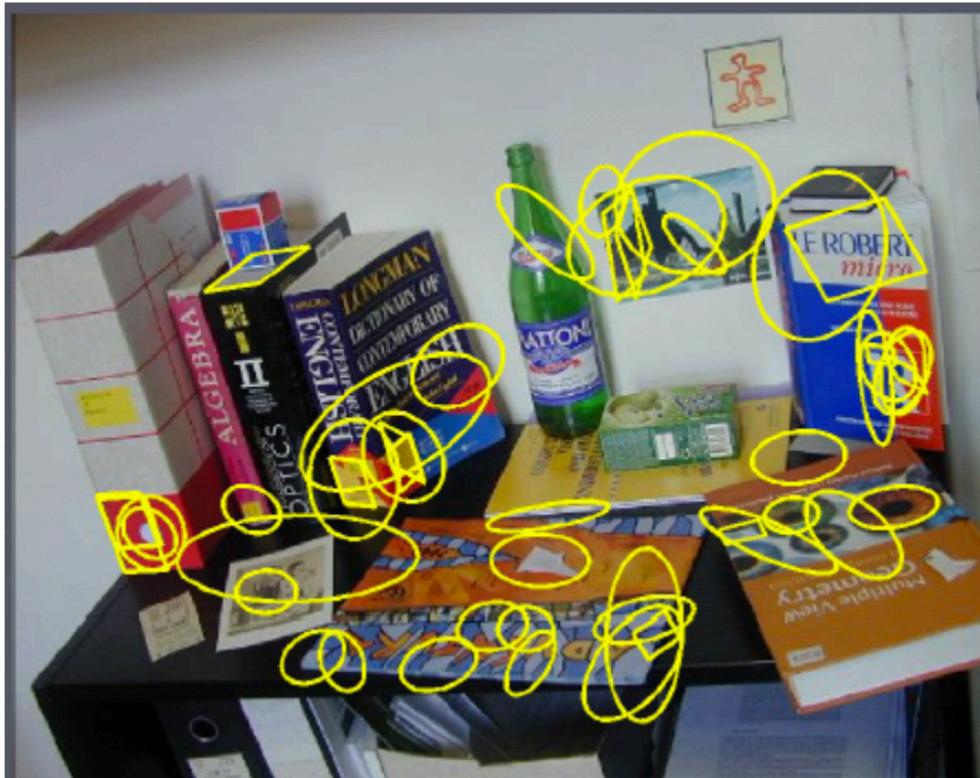
Automatic mosaicing



<http://matthewwalunbrown.com/autostitch/autostitch.html>

# Applications

Wide baseline stereo



[Image from T. Tuytelaars ECCV 2006 tutorial]